


page table

page # \leftrightarrow frame #

logical address



This 'chopping' operation is done automatically by MMU.


of bits used to represent the offset is determined by the page size

page size (4Kb) would need 12 bits

$$2^{12} = 4K$$

Why we need to use hierarchical structure?
 \uparrow for page table

DB design's perspective




↑ for page table

Size of page table varies

Solution: store page table into pages.


e.g. process has n pages w/ page size N Bytes
for user process data/instr



page table is stored and managed by kernel.

it could use different sizes of the page.

for kernel data structures.



32 bits ^{ways} page size is 4KB \Rightarrow 12 bits for offset

$$(2^{12} = 4K)$$

if we store page table
into pages (4KB size)

then

One page (4KB) would
store 1K entries

[P.T.E needs 4 Bytes)

So # of bits to represent the
"address" of each entry

is 10 bits
 $2^{10} = 1K$

register	access	the	a
Cache	"	"	b
mem	"	"	c

A two level page table, given a logical

address a_1 , the reference to a_1 would need _____ time units

a) a

b) $a+b$

c) $3 \cdot c$

d) $4 \cdot c$

from effective memory access perspective,

Contiguous allocation is the most effective.

effective access time is $(2)a + c$

OS maintains a base [register] & length [register] for each process

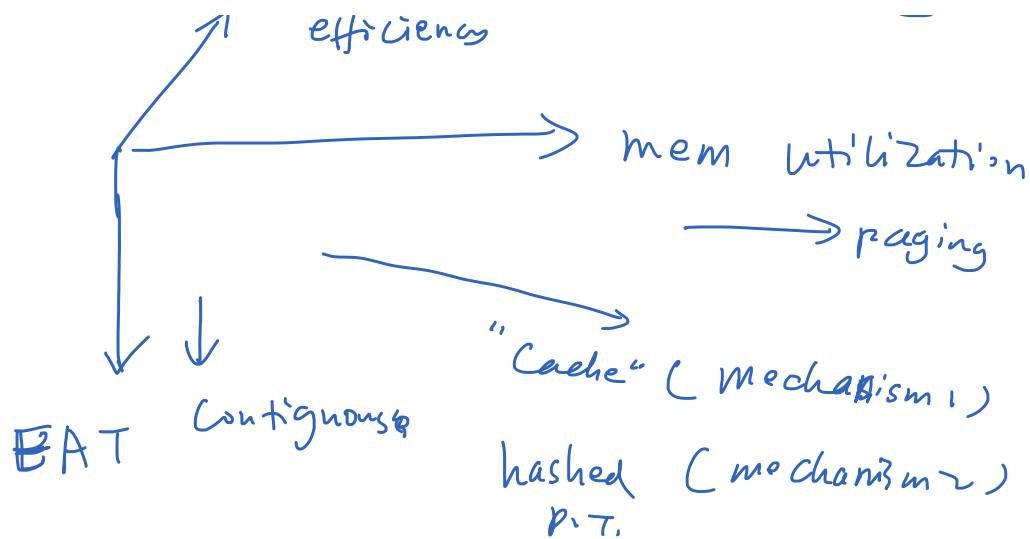
Given a logical address a_1 ,

physical address is $\text{base} + a_1$

EAT = access to base register
length

$2a + c$

management efficiency



P#	f#

if we could combine all the page tables together, then ~~that~~ we could find that P# values are not unique.

The other way around,

if there is 4GB mem (physical)
page size is 4KB then

$$\text{there is } \frac{4GB}{4KB} = \underline{\underline{1M \text{ frames}}}$$

assume P.T.E uses 4 Bytes,
then an inverted page table needs
 $1M * 4 \text{ Bytes} = \underline{\underline{4MB}}$

$$1M * 4 \text{ Bytes} = \underline{\underline{4M \text{ B}}}$$

[illegible]

data structure maintained by
D.S. for both
used and
free memory

Compared to the previous
page table mechanism

dynamically manage x
page tables
 Q
 a bit map

Current # of processes in the system

V.M.

if we have 4GB physical memory
(main)

On the Average, a process needs
100MB mem in total.

m.m. could store $\frac{4 \text{ Kbs}}{100 \mu\text{s}} \doteq 40.9 \text{ processes}$

if we load only 10MB for each process.


then m.m could store

$$\frac{4\text{TB}}{10\text{MB}} = 409 \text{ processes}$$

by using VM, these 409 processes

have an illusion of m.m. size
of

$$409 \times 100\text{MB} = \underline{\underline{40\text{GB}}}$$



Total
 logical
 address space
 >> 4GB
 physical
 address space

```

if ( ) {
  if-block
} else {
  else-
  block
}
  
```

load to entire process into
 m.m, OS allocates
 2 pages

if we go up demand paging

else-
block
}

if we go w/ demand paging
then
1 at most.

Page fault rate p

$p=0 \Rightarrow$ no page fault (all the
needed pages are
in mem)

$p=1 \Rightarrow$ every reference causes
a page fault.

page size is 1 byte.

4KB page size

each instruction ~~is~~ needs 4 Bytes.

One page may have 1K instructions

each instruction involves one address

one page may have 1K mem
references

for (i < 100) {

x

}

100 references
to x

2K for each page

then 400,000 references

means $\frac{400,000}{2K} = 200 \text{ pages}$

$200 \times 4KB = \underline{\underline{800 \text{ Kbytes}}}$

page 18.

if page D of process 2 is replaced.

page table for p₂

then

6	V
	i
	i
7	V

page D is swapped out
to disk

page table for p₁ is now

3	✓
4	✓
5	✓
2	✓

