

# **CSI 4500 Operating Systems**

## **Lecture 2**

### **Processes and Threads**

# Today's Goals

## ■ What are Processes?

- Process Concept
- Process Scheduling
- Operations on Processes

## ■ Understanding Threads

- Thread Dispatching
- Beginnings of Thread Scheduling

# Process Concept

## ■ Process – a program in execution

- Operating system abstraction to represent what is need to run a program

## ■ A process includes:

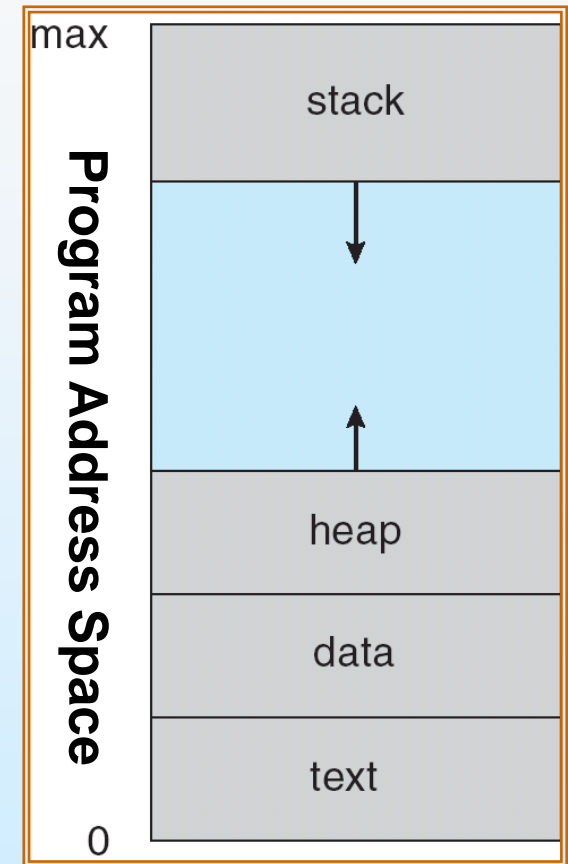
- Sequential Program Execution Instruction Stream
  - ▶ Code executed as a *single, sequential* stream of execution
  - ▶ Includes State of CPU registers
- Protected Resources:
  - ▶ Main Memory State (contents of Address Space)
  - ▶ I/O state (i.e. file descriptors)

## ■ Processes can be described as either:

- **I/O-bound process**
  - ▶ spends more time doing I/O than computations
  - ▶ many short CPU bursts
- **CPU-bound process**
  - ▶ spends more time doing computations
  - ▶ few very long CPU bursts


# Process in Memory

- Address space  $\Rightarrow$  the set of accessible addresses + state associated with them:
  - For a 32-bit processor there are  $2^{32}$  bits addresses
- Read or write operation to an address?
  - Regular memory access
  - Exception
  - I/O operation

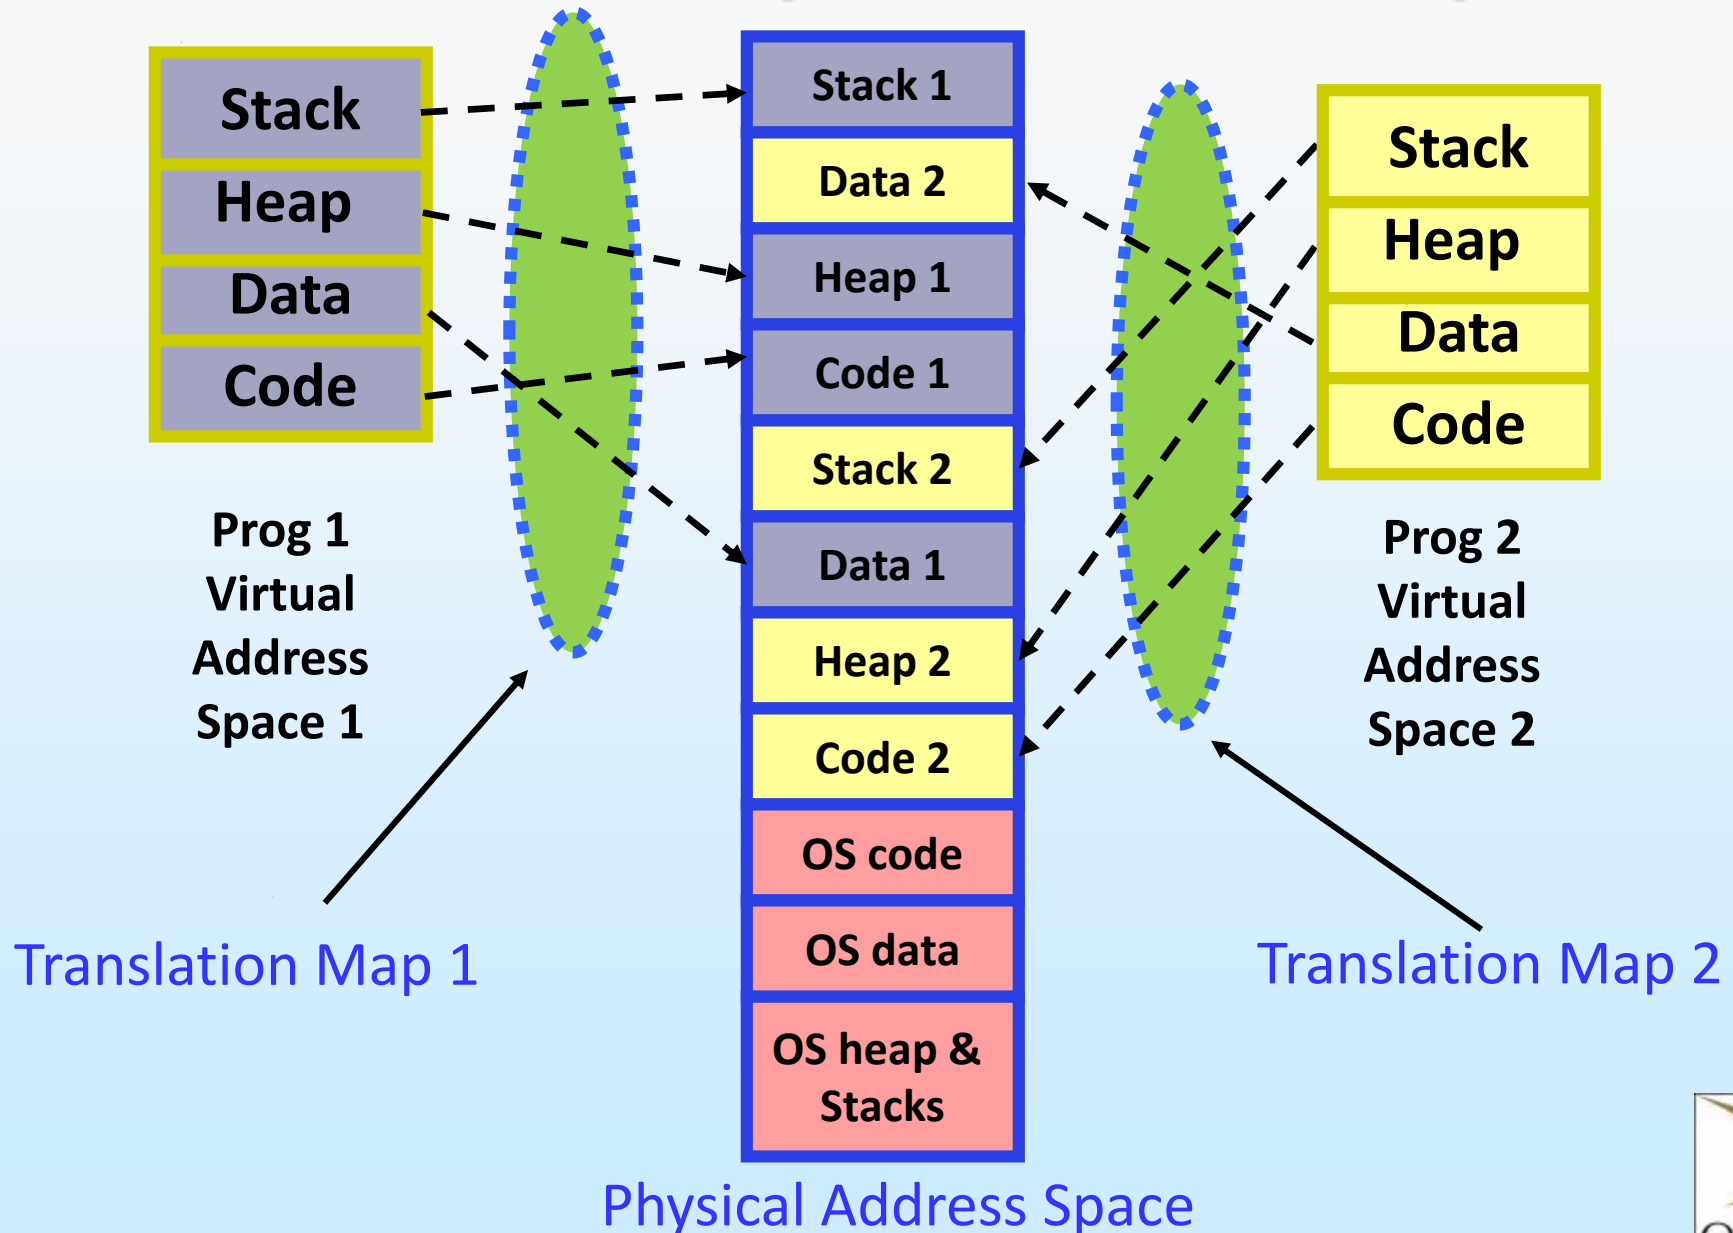


# Memory Layout of a Process

- **Text segment** contains the machine-language instructions of the program run by a process.
- **Initialized data segment** contains global and static variables that explicitly initialized.
- **Uninitialized data segment** contains global and static variables that not explicitly initialized
- **Stack** contains stack frames. One stack frame is allocated for each currently called function. A frame stores the function's local variables (automatic variables), arguments, and return value.
- **Heap** is an area of memory can be dynamically allocated at run time.

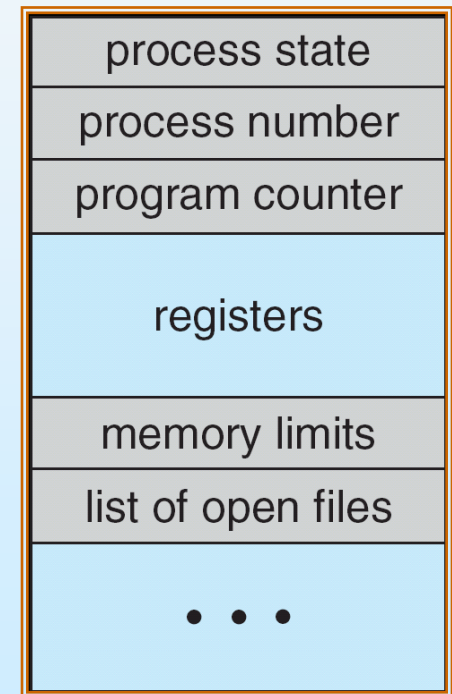


# Illustration of Separate Address Space



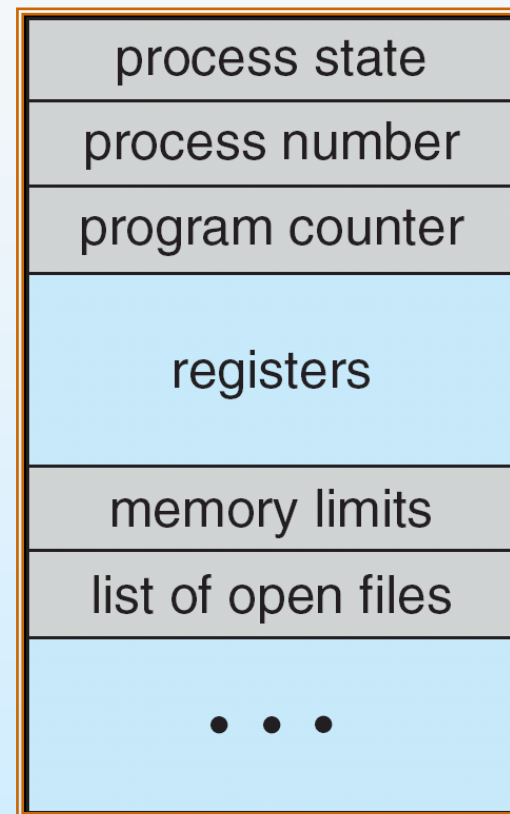
# Process Control Block (PCB)

- The current state of process held in a process control block (PCB):
  - This is a “snapshot” of the execution and protection environment
  - Information associated with each process
    - ▶ Process state
    - ▶ Program counter
    - ▶ CPU registers
    - ▶ CPU scheduling information
    - ▶ Memory-management information
    - ▶ Accounting information
    - ▶ I/O status information



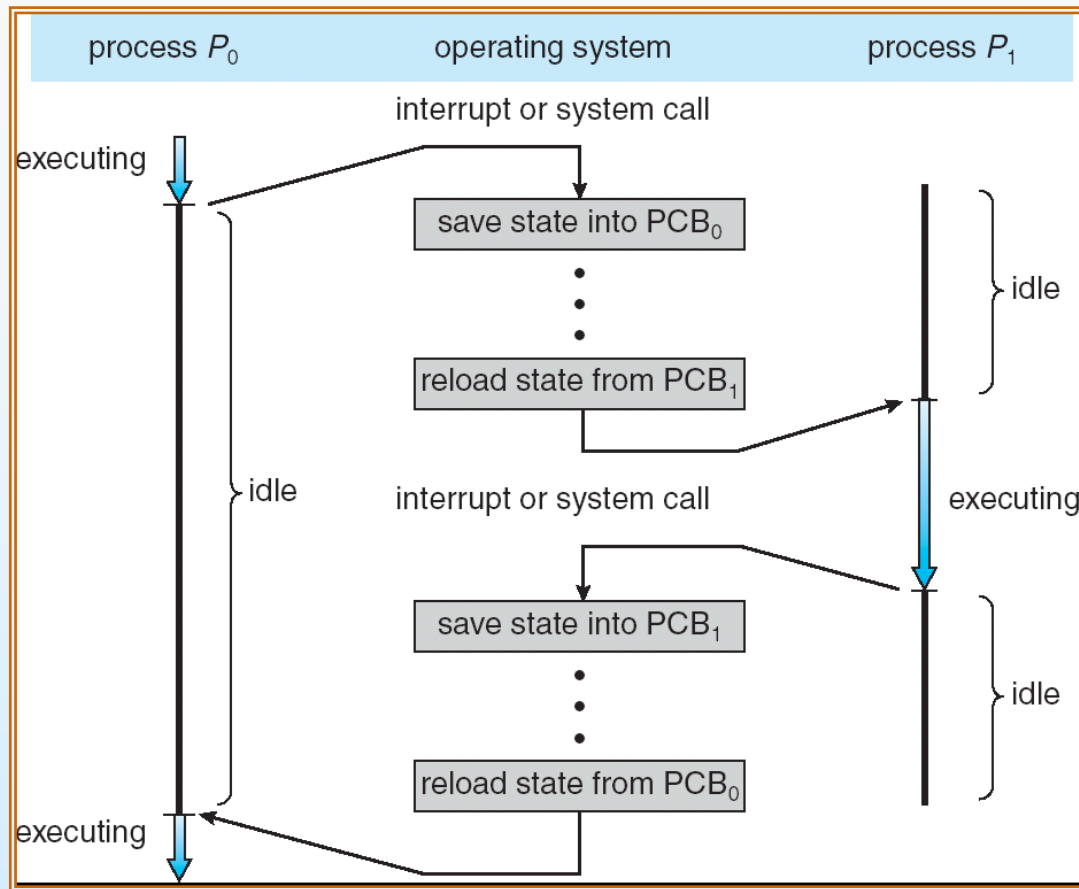
# How do we multiplex processes?

- Process control block (PCB):
  - Only one PCB active at a time
- Assign CPU time to different processes (**Scheduling**):
  - Only one process “running” at a time
  - Give more time to important processes (**Priority**)
- Give pieces of resources to different processes (**Protection**):
  - Controlled access to non-CPU resources
  - Sample mechanisms:
    - ▶ Memory Mapping: Give each process their own address space
    - ▶ Kernel/User duality: Arbitrary multiplexing of I/O through system calls





# CPU Switch From Process to Process

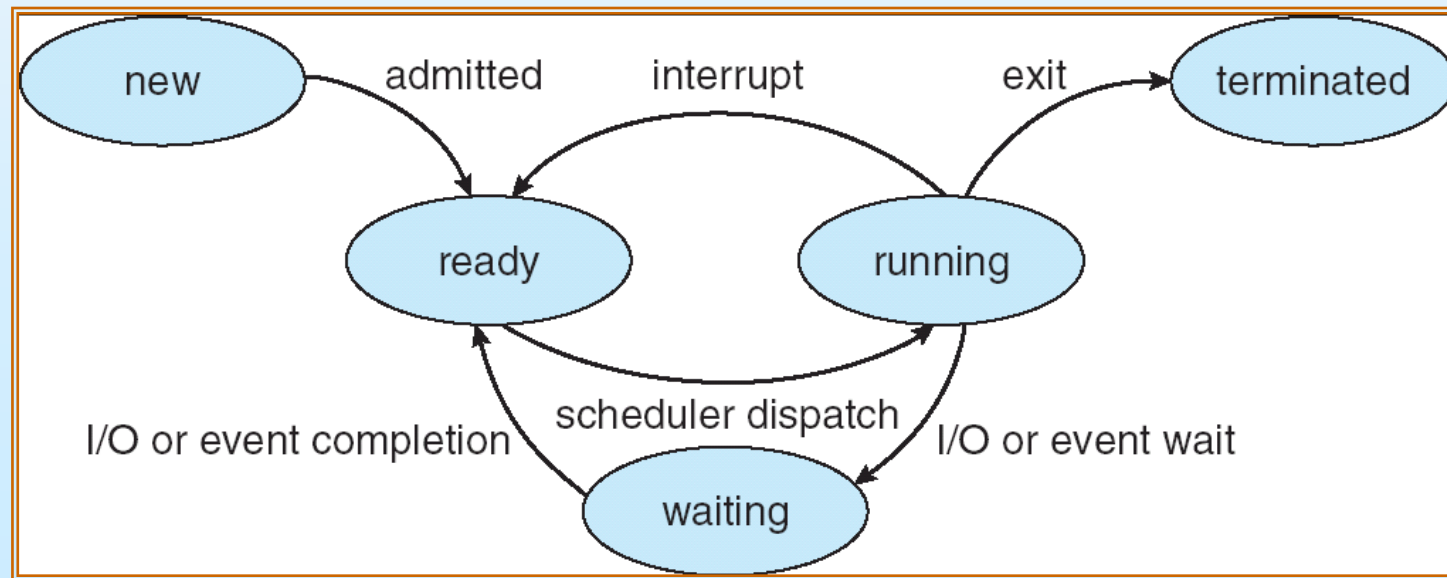


- This is also called a “context switch”
- During context switch system does no useful work
  - Overhead dependent on hardware support
  - What are they?

# Process State

■ As a process executes, it changes *state*

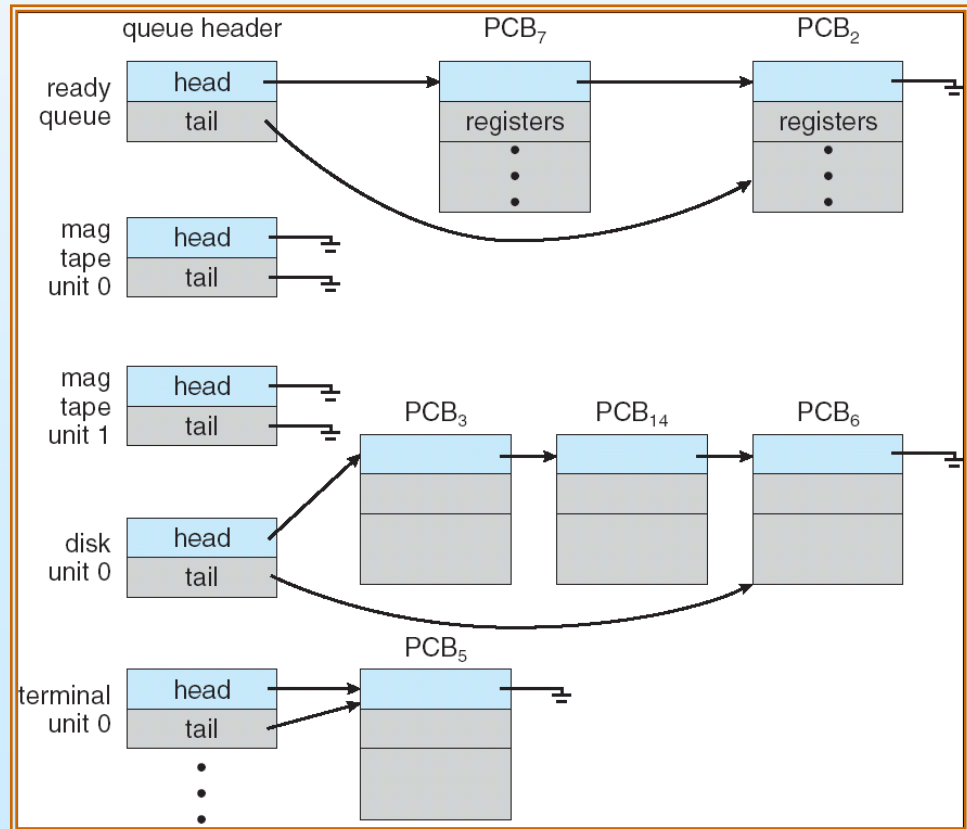
- **new**: The process is being created
- **ready**: The process is waiting to run
- **running**: Instructions are being executed
- **waiting**: Process waiting for some event to occur
- **terminated**: The process has finished execution



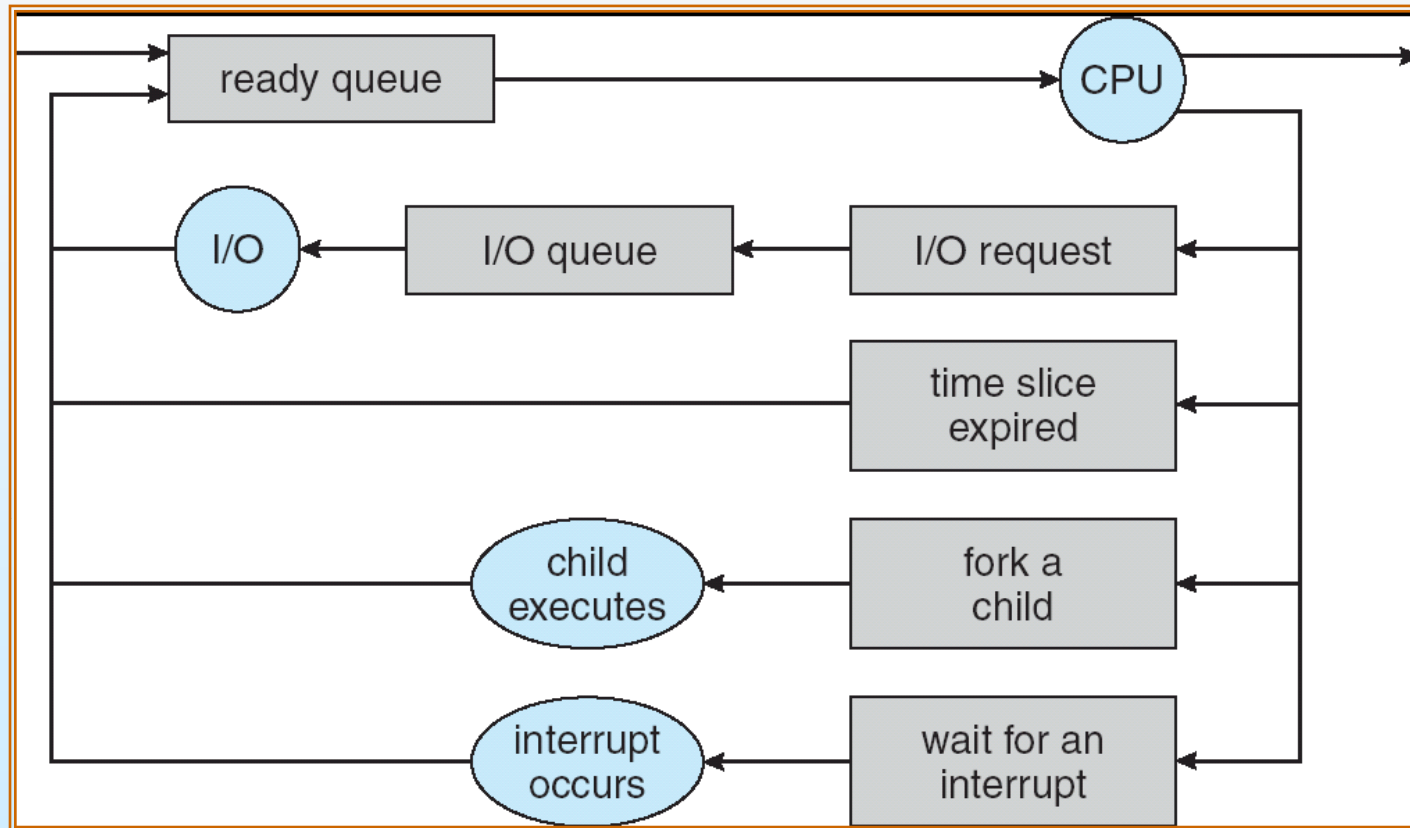
# Process Scheduling Queues

- **Job queue** – set of all processes in the system
- **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
- **Device queues** – set of processes waiting for an I/O device
- **Scheduling**

Processes (PCBs)  
migrate among the  
various queues

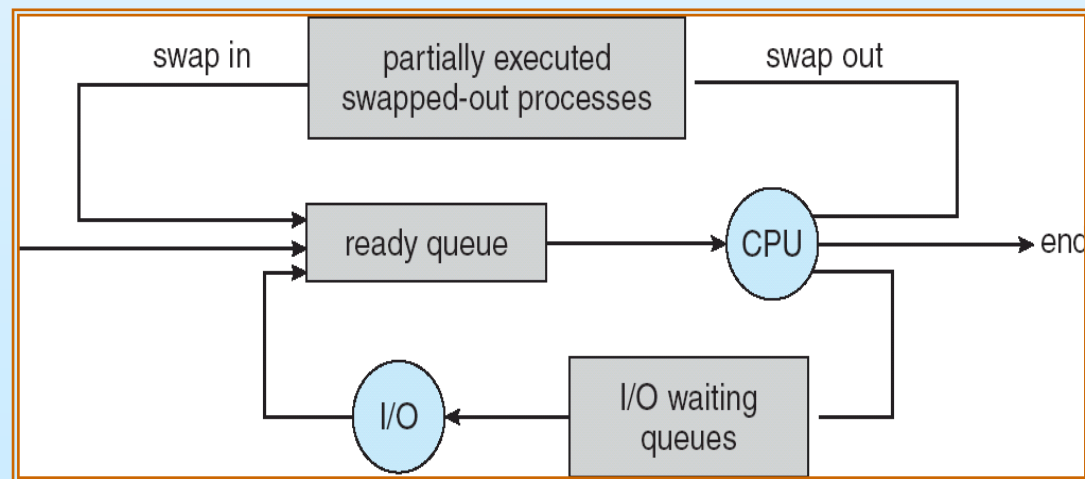


# Representation of Process Scheduling



# Schedulers

- **Long-term scheduler** (or job scheduler) – selects which processes should be brought into the ready queue
  - Long-term scheduler is invoked very infrequently
  - The long-term scheduler controls the *degree of multiprogramming*
- **Short-term scheduler** (or CPU scheduler) – selects which process should be executed next and allocates CPU
  - Short-term scheduler is invoked very frequently (milliseconds)  $\Rightarrow$  (must be fast)
- **Medium-term scheduler**




# How to Create a Process?

- Parent process create children processes, which, in turn create other processes, forming a tree of processes
- Must construct new PCB
- Resource sharing strategies
  - Parent and children share all resources (I/O states, address space information)
  - Children share subset of parent's resources
  - Parent and child share no resources (Unix *exec()*)
- Execution
  - Parent and children execute concurrently
  - Parent waits until children terminate

# Process Creation – C example

```
int main() {
    pid_t  pid;
    pid = fork(); /* fork another process */
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait (NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```



# Process Termination

- Process executes last statement and asks the operating system to delete it (**exit**)
  - Output data from child to parent (via **wait**)
  - Process' resources are deallocated by operating system
- Parent may terminate execution of children processes (**abort**)
  - Child has exceeded allocated resources
  - Task assigned to child is no longer required
  - If parent is exiting
    - ▶ Some operating systems do not allow child to continue if its parent terminates
      - All children terminated - *cascading termination*



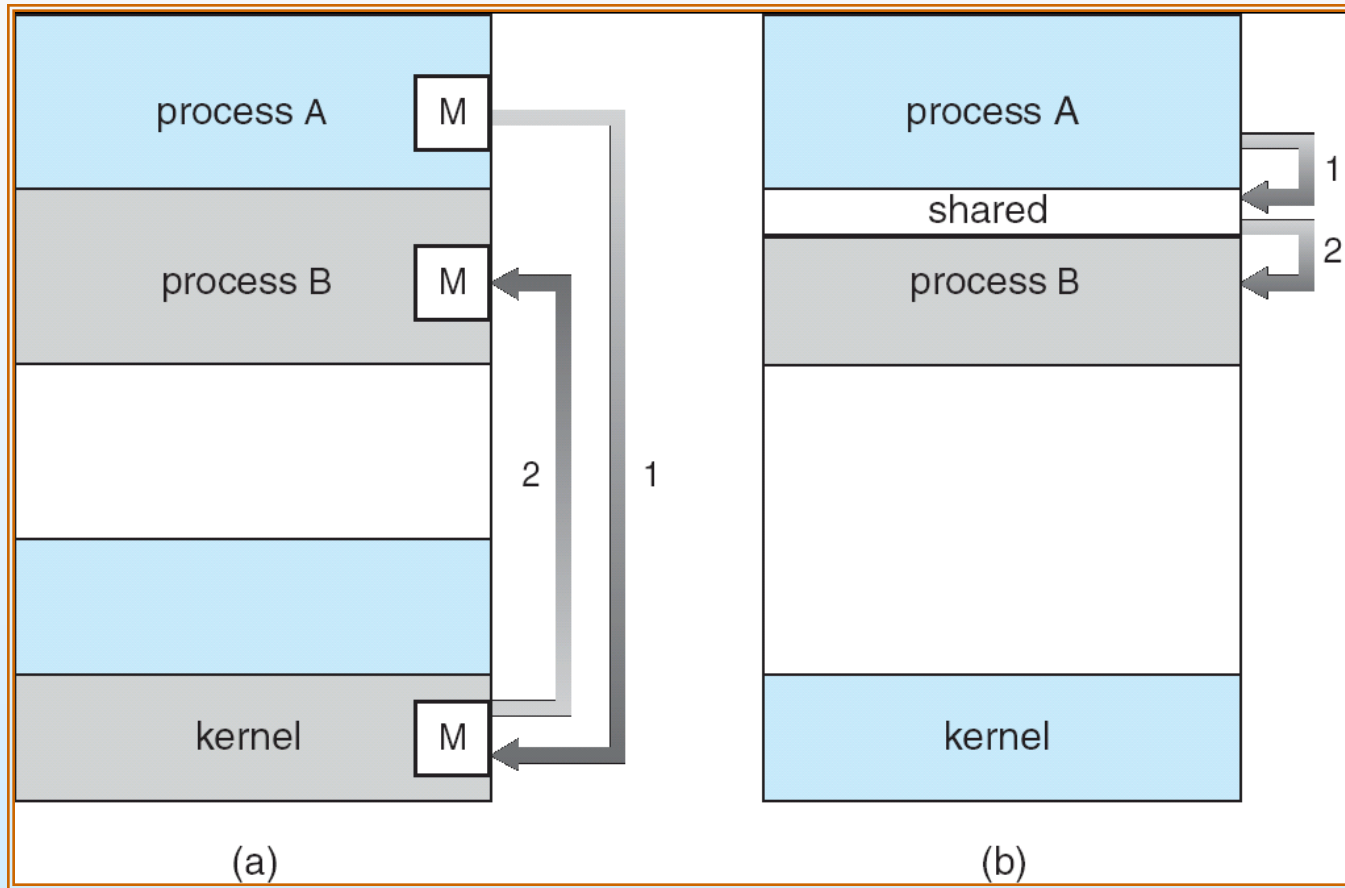
# Cooperating Processes

- **Independent** process cannot affect or be affected by the execution of another process
- **Cooperating** process can affect or be affected by the execution of another process
- Advantages of process cooperation
  - Information sharing
  - Computation speed-up
  - Modularity
  - Convenience
- Disadvantages?
  - High Creation/memory Overhead
  - (Relatively) High Context-Switch Overhead

# Interprocess Communication (IPC)

- Mechanism for processes to communicate and to synchronize their actions
- **Message system** — processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
  - **send**(*message*) – message size fixed or variable
  - **receive**(*message*)
- If  $P$  and  $Q$  wish to communicate, they need to:
  - establish a *communication link* between them
  - exchange messages via send/receive
- Implementation of communication link
  - physical (e.g., shared memory, hardware bus)
  - logical (e.g., logical properties)

# Communications Models



(a) Message Passing

(b) Shared-Memory Mapping

# Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?

# Direct Communication

## ■ Processes must name each other explicitly:

- **send** ( $P, message$ ) – send a message to process  $P$
- **receive**( $Q, message$ ) – receive a message from process  $Q$

## ■ Properties of communication link

- Links are established automatically
- A link is associated with exactly one pair of communicating processes
- Between each pair there exists exactly one link
- The link may be unidirectional, but is usually bi-directional

# Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports)
  - Each mailbox has a unique id
  - Processes can communicate only if they share a mailbox
- Properties of communication link
  - Link established only if processes share a common mailbox
  - A link may be associated with many processes
  - Each pair of processes may share several communication links
  - Link may be unidirectional or bi-directional
- Operations
  - create a new mailbox
  - send and receive messages through mailbox
  - destroy a mailbox
- Primitives are defined as:
  - **send**(*A, message*) – send a message to mailbox A
  - **receive**(*A, message*) – receive a message from mailbox A

# Indirect Communication

## ■ Mailbox sharing

- $P_1$ ,  $P_2$ , and  $P_3$  share mailbox A
- $P_1$  sends;  $P_2$  and  $P_3$  receive
- Who gets the message?

## ■ Solutions

- Allow a link to be associated with at most two processes
- Allow only one process at a time to execute a receive operation
- Allow the system to select arbitrarily the receiver.
  - ▶ Sender is notified who the receiver was.

# Synchronization

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
  - **Blocking send** has the sender block until the message is received
  - **Blocking receive** has the receiver block until a message is available
- **Non-blocking** is considered **asynchronous**
  - **Non-blocking send** has the sender send the message and continue
  - **Non-blocking receive** has the receiver receive a valid message or null



# Buffering

- Queue of messages attached to the link
  - Zero capacity – 0 messages
    - ▶ *Sender must wait for receiver (rendezvous)*
  - Bounded capacity – finite length of  $n$  messages
    - ▶ *Sender must wait if link full*
  - Unbounded capacity – infinite length
    - ▶ *Sender never waits*

# Bounded-Buffer – Shared Memory Solution

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

**Share Data**

```
while (true) {
    /* Produce an item */
    while (((in = (in + 1) % BUFFER_SIZE count) == out)
        ; /* do nothing -- no free buffers */
    buffer[in] = item;
    in = (in + 1) % BUFFER_SIZE;
}
```

**Insert()**

```
while (true) {
    while (in == out)
        ; // do nothing -- nothing to consume
    item = buffer[out]; // remove an item from the buffer
    out = (out + 1) % BUFFER_SIZE;
    return item;
}
```

**Remove()**



# Process Summary

## ■ Processes have two parts

- Sequence of Execution Stream
- Resources

## ■ Concurrency accomplished by multiplexing CPU Time:

- Unloading current thread (PC, registers)
- Loading new thread (PC, registers)
- Such context switching may be voluntary (`yield()`, I/O operations) or involuntary (timer, other interrupts)

## ■ Protection accomplished restricting access:

- Memory mapping isolates processes from each other
- Dual-mode for isolating I/O, other resources

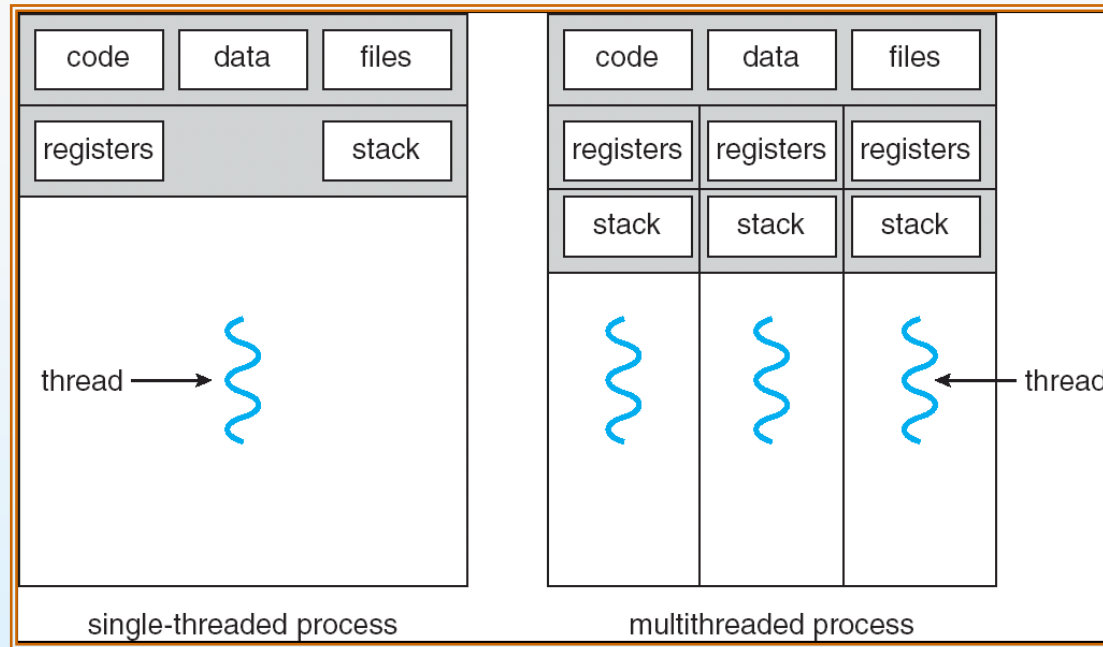
## ■ Cooperating of Processes

- Shared Memory Communication
- Message Communication

# Multiple Threads within a Process

- Process: Operating system abstraction to represent what is needed to run a single (**multithreaded**) program
- Two parts:
  - Multiple Threads
    - ▶ Each thread is a *single, sequential stream of execution*
  - Protected Resources:
    - ▶ Main Memory State (contents of Address Space)
    - ▶ I/O state (i.e. file descriptors)
- Why separate the concept of a thread from that of a process?
  - Heavyweight Process  $\equiv$  Process with only one thread

# Single and Multithreaded Processes



## ■ Threads encapsulate concurrency

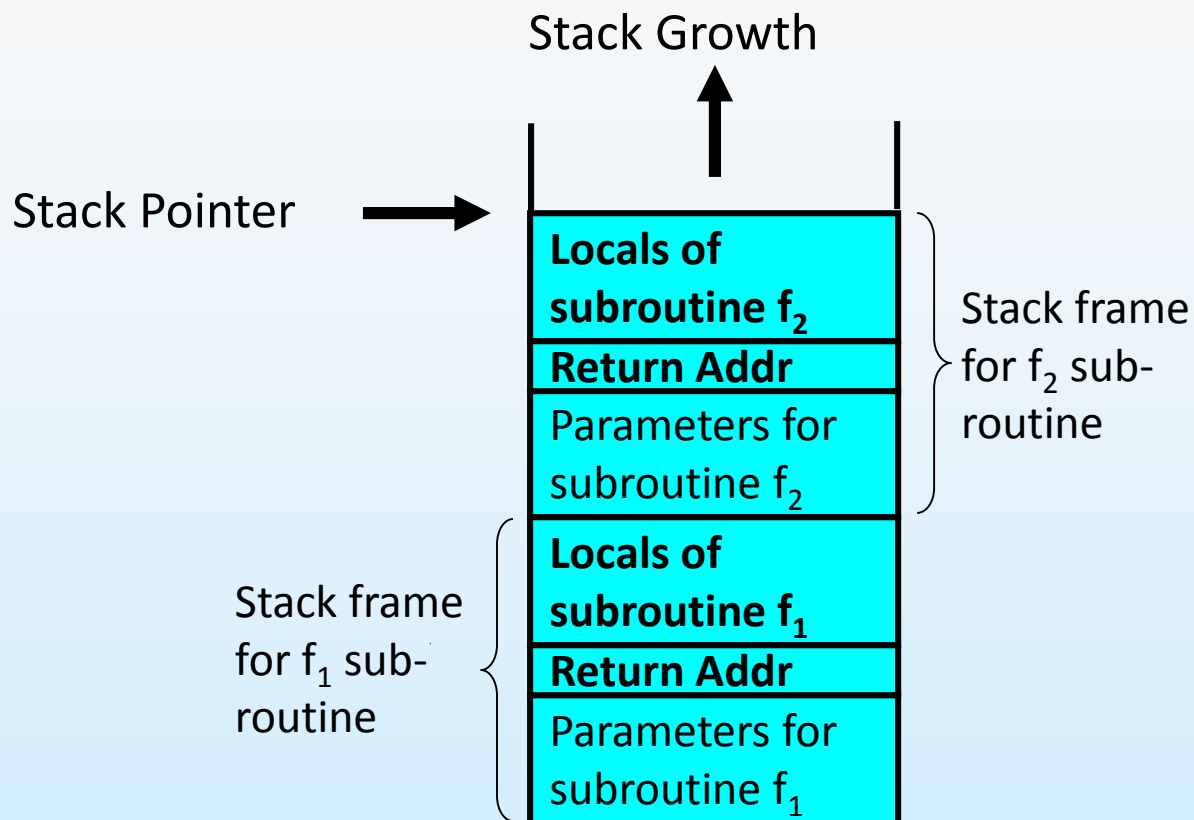
- “Control” component of a process

## ■ Address spaces encapsulate protection

- “Passive” component of a process

# Recall: Call Stack Example

```
f1(int tmp) {  
    ...  
    f2(a, b);  
    ...  
}  
  
f2(int a,  
float b) {  
    int i;  
    ...  
}
```



- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

# Single-Threaded Example

- Imagine the following Pseudo program:

```
main() {  
    ...  
    PrintDigitPI("pi.txt");  
    PrintDigitE("e.txt");  
    ...  
}
```

- What is the behavior here?
  - Program would never print out **E** (mathematical constant)
  - Why?
    - ▶ `PrintDigitPI()` would never finish

# Use of Threads

## ■ Version of program with Threads:


```
main() {  
    ...  
    CreateThread(PrintDigitPI("pi.txt"));  
    CreateThread(PrintDigitE("e.text"));  
    ...  
}
```

## ■ What does “CreateThread()” do?

- Start independent thread running given procedure

## ■ What is the program behavior now?


- Now, you would actually see the both **PI** and **E**
- This *should* behave as if there are two separate CPUs





# Memory View of Two Threads

- If we stopped this program and examined it with a debugger, we would see
  - Two sets of CPU registers
  - Two sets of Stacks
- Questions:
  - How do stacks locate relative to each other?
  - What maximum size for the stacks?
  - What happens if threads violate?
  - How might you catch violations?

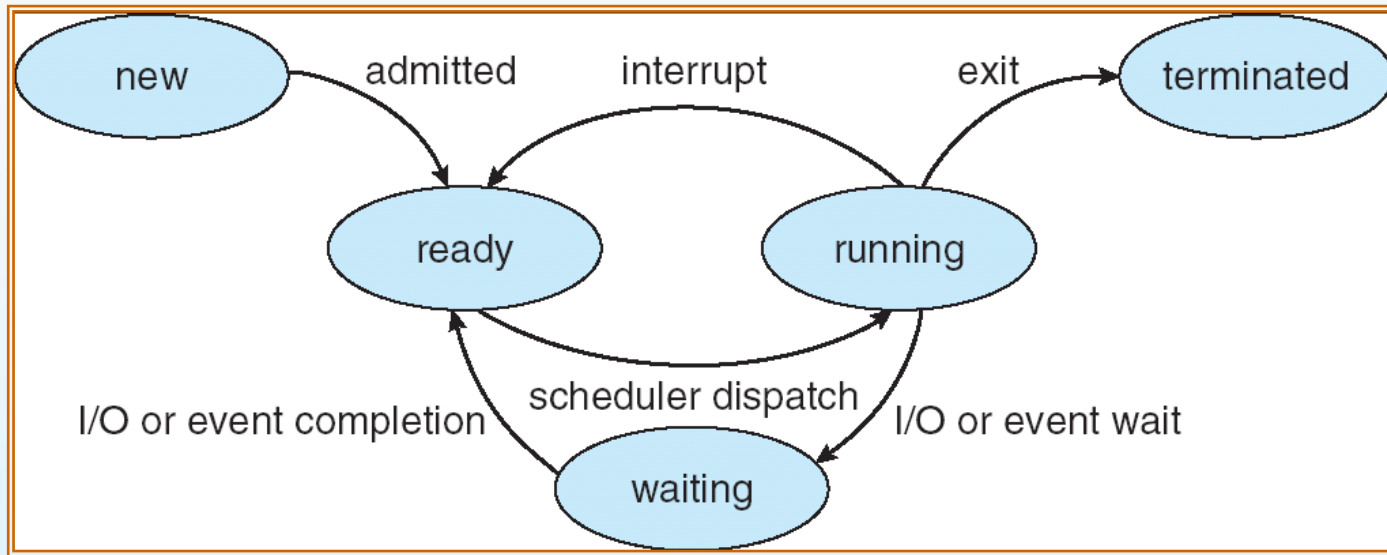


Address Space

# Thread State

- Each Thread has a *Thread Control Block* (TCB)
  - **Execution State**: CPU registers, program counter, pointer to stack
  - **Scheduling info**: State, priority, CPU time
  - Various Pointers (for implementing scheduling queues)
  - Etc.
- In Java: “Thread” is a class that includes the TCB
- OS Keeps track of TCBs in protected memory
  - In Array, or Linked List, or ...

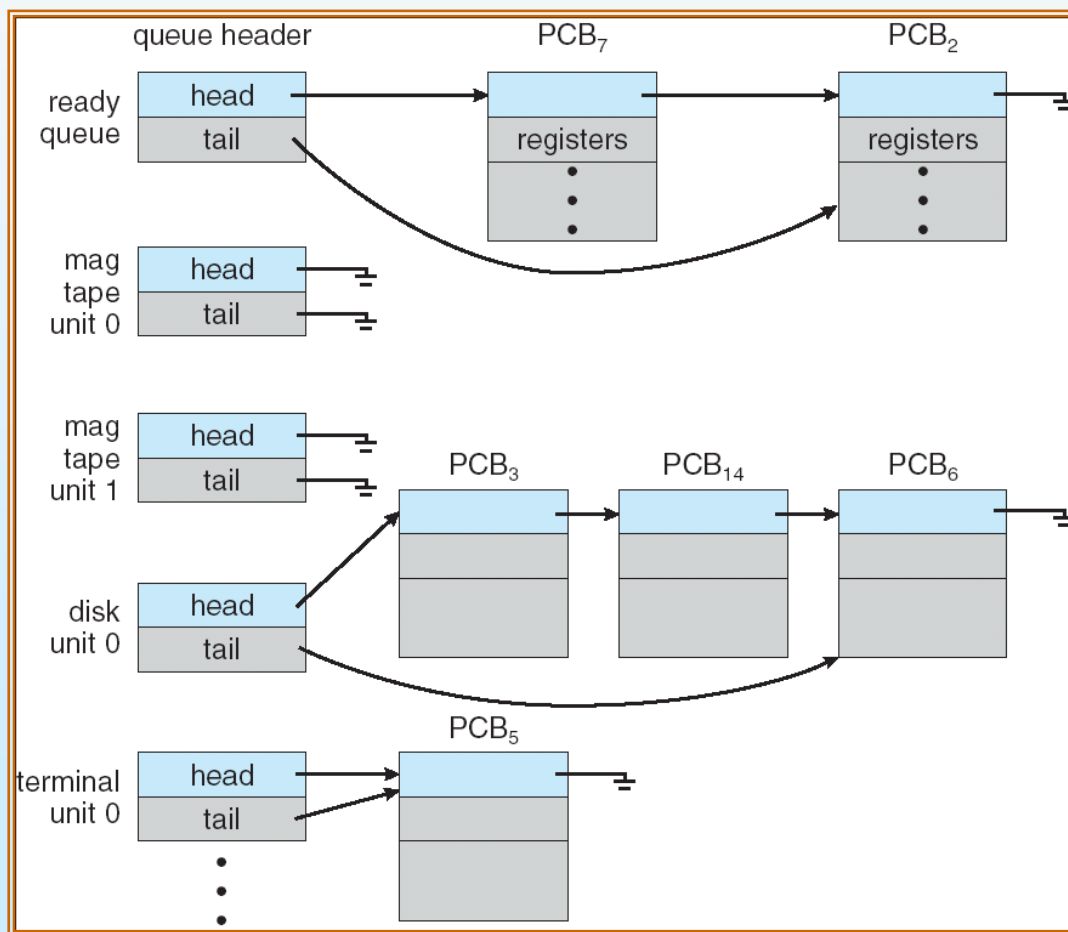
# Recall: Lifecycle of a Thread (or Process)



- As a thread executes, it changes state:
  - **new**: The thread is being created
  - **ready**: The thread is waiting to run
  - **running**: Instructions are being executed
  - **waiting**: Thread waiting for some event to occur
  - **terminated**: The thread has finished execution
- “Active” threads are represented by their TCBs
  - TCBs organized into queues based on their state

# Ready Queue And Various I/O Device Queues

- Thread not running  $\Rightarrow$  TCB is in some other scheduler queues
  - Queues exist based on device/signal/condition
  - Each queue may have a different scheduler policy



# OS operates flow

## ■ Conceptual view of the operating system

```
Loop {  
    RunThread();  
    ChooseNextThread();  
    SavecurrentTCB();  
    LoadStateOfCPU(newTCB);  
}
```

## ■ This is an *infinite* loop

- One could argue that this is all that the OS does

## ■ When we ever exit this loop???

# Running a thread

## ■ How do I run a thread?

- Load its state (registers, PC, stack pointer) into CPU
- Load environment (virtual memory space, etc)
- Jump to the PC

## ■ How does the dispatcher get control back?

- **Internal events**: thread returns control voluntarily
- **External events**: thread gets *preempted*

# Internal Events

- Blocking on I/O
  - The action of requesting I/O implicitly yields the CPU
- Waiting on a “signal” from other thread
  - Thread asks to wait and thus yields the CPU
- Thread executes a `yield()`
  - Thread volunteers to give up CPU

```
PrintDigitPI() {  
    while(TRUE) {  
        PrintNextDigit();  
        yield();  
    }  
}
```

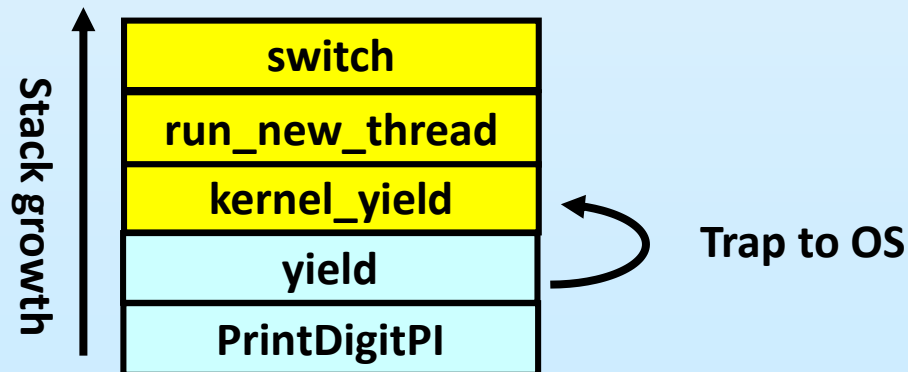
# Stack for Yielding Thread

## ■ How do we run a new thread?

```
run_new_thread() {  
    newThread = PickNewThread();  
    switch(curThread, newThread);  
    ...  
}
```

## ■ How does dispatcher switch to a new thread?

- Save anything next thread may trash: PC, regs, stack
- Maintain isolation for each thread

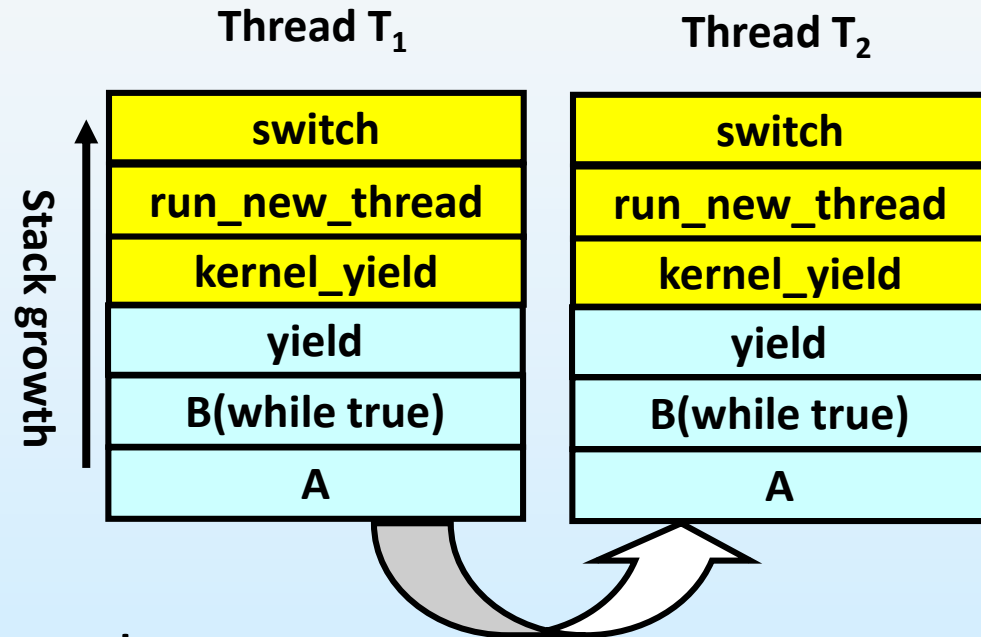




# What do the stacks look like?

- Consider the following code blocks:

```
proc A() {  
    B();  
}  
proc B() {  
    while(TRUE) {  
        yield();  
    }  
}
```

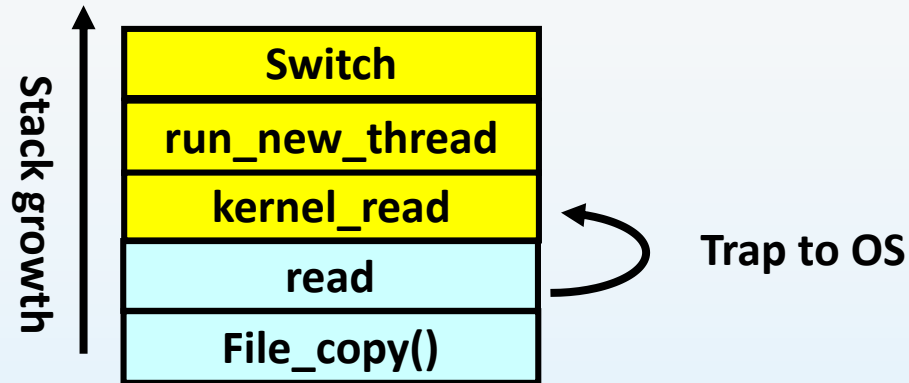


- Suppose we have two threads:
  - Threads T<sub>1</sub> and T<sub>2</sub>

# Saving/Restoring state

```
Switch(curThread,newThread) {  
    /* Unload old thread */  
    TCB[curThread].regs.r7 = CPU.r7;  
    ...  
    TCB[curThread].regs.r0 = CPU.r0;  
    TCB[curThread].regs.sp = CPU.sp;  
    TCB[curThread].regs.retpc = CPU.retpc; /*return addr*/  
  
    /* Load and execute new thread */  
    CPU.r7 = TCB[newThread].regs.r7;  
    ...  
    CPU.r0 = TCB[newThread].regs.r0;  
    CPU.sp = TCB[newThread].regs.sp;  
    CPU.retpc = TCB[newThread].regs.retpc;  
    return; /* Return to CPU.retpc */  
}
```

# Thread blocks on I/O



## ■ What happens when a thread requests a block of data from the file system?

- User code invokes a system call
- Read operation is initiated
- Run new thread/switch

## ■ Thread communication similar

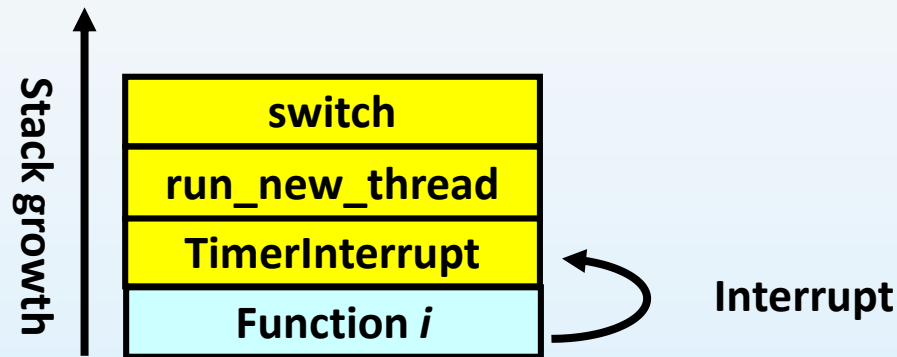
- Wait for Signal/Join
- Networking

# External Events

- What happens if thread never does any I/O, never waits, and never yields control?
  - Could the *PrintDigitPI* program grab all resources and never release the processor?
    - ▶ What if it didn't print anything?
  - Must find way that dispatcher can regain control!
- Answer: Utilize External Events
  - **Interrupts**: signals from hardware or software that stop the running code and jump to kernel
  - **Timer**: like an alarm clock that goes off every some many milliseconds
- If we make sure that external events occur frequently enough, can ensure dispatcher runs

# To Acquire Control via Interrupt

- Solution to our dispatcher problem
  - Use the timer interrupt to force scheduling decisions



- Timer Interrupt routine:

```
TimerInterrupt() {  
    DoPeriodicChecking();  
    run_new_thread();  
}
```

- I/O interrupt: same as timer interrupt except that `DoPeriodicChecking()` replaced by `ServiceIO()`.

# Choosing a Thread to Run

## ■ How does Dispatcher decide what to run next?

- Zero ready threads – dispatcher loops
  - ▶ Alternative is to create an “idle thread”
  - ▶ Can put machine into low-power mode
- Exactly one ready thread – easy
- More than one ready thread: scheduling

## ■ Possible priorities:

- LIFO (last in, first out):
  - ▶ put ready threads on front of list, remove from front
- FIFO (first in, first out):
  - ▶ Put ready threads on the tail of list, pick them from front
  - ▶ Fair policy
- Priority queue:
  - ▶ keep ready list sorted by TCB priority field
- Pick one at random

# Threads Summary

- The state of a thread is contained in the TCB
  - Registers, PC, stack pointer
  - States: New, Ready, Running, Waiting, or Terminated
- Multithreading provides simple illusion of multiple CPUs
  - Switch registers and stack to dispatch new thread
  - Provide mechanism to ensure dispatcher regains control
- Switch routine
  - Can be very expensive if many registers
  - Must be very carefully constructed!
- Many scheduling options
  - Decision of which thread to run complex enough for complete lecture