

CSE 4500: Operating Systems

Lecture 3

Process/Thread Synchronization

Background

- Parallelism can provide a distinct way of conceptualizing problems.
- Concurrent access to shared data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers.
 - We can do so by having an integer **count** that keeps track of the number of full buffers.
 - Initially, **count** is set to 0.
 - It is incremented by the producer after it produces a new buffer
 - It is decremented by the consumer after it consumes a buffer.

Producer and Consumer

```
while (true) {  
    /* produce an item and put in nextProduced */  
    while (count == BUFFER_SIZE)  
        ; // do nothing  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

```
while (true) {  
    while (count == 0)  
        ; // do nothing  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
    /* consume the item in nextConsumed  
}
```

Race Condition

- `count++` could be implemented as
 - `register1 = count`
 - `register1 = register1 + 1`
 - `count = register1`
- `count--` could be implemented as
 - `register2 = count`
 - `register2 = register2 - 1`
 - `count = register2`
- Consider this execution interleaving with “count = 5” initially:

S0: producer execute <code>register1 = count</code>	{register1 = 5}
S1: producer execute <code>register1 = register1 + 1</code>	{register1 = 6}
S2: consumer execute <code>register2 = count</code>	{register2 = 5}
S3: consumer execute <code>register2 = register2 - 1</code>	{register2 = 4}
S4: producer execute <code>count = register1</code>	{count = 6}
S5: consumer execute <code>count = register2</code>	{count = 4}
- Race conditions can occur when operations on shared variables are not ***atomic***.

Definitions

- **Synchronization**: using atomic operations to ensure cooperation between threads
- **Critical Section**: piece of code that only one thread can execute at once. Only one thread at a time will get into this section of code.
 - **Mutual Exclusion**: ensuring that only one thread does a particular thing at a time
 - **Progress**: selecting a thread to enter cannot postpone indefinitely
 - **Bounded waiting**: before entering into the critical section
- **Lock**: prevents someone from doing something
 - Lock before accessing shared data
 - Unlock after accessing shared data
 - Wait if locked

Important idea: all synchronization involves waiting

Where to synchronize?

Programs	Shared Programs
API	Locks Semaphores Monitors Send/Receive
Hardware	Load/Store Disable Ints Test&Set Comp&Swap

- We are going to implement various higher-level synchronization primitives
 - Need to provide primitives useful at user-level

Peterson's Solution

- A solution for two processes
- Assume that the LOAD and STORE instructions are atomic.
 - atomic == cannot be interrupted.
- The two processes share two variables:
 - int **turn**;
 - Boolean **flag[2]**
- The variable **turn** indicates whose turn it is to enter the critical section.
- The **flag** array is used to indicate if a process is ready to enter the critical section.
 - **flag[i] = true** implies that process P_i is ready!

Algorithm for Process P_i

```
while (true) {  
    flag[i] = TRUE;  
    turn = j;  
    while ( flag[j] && turn == j) //busy waiting  
        ;  
    CRITICAL SECTION  
    flag[i] = FALSE;  
    REMAINDER SECTION  
}
```


Synchronization Hardware

- Many systems provide hardware support for critical section code
- Uniprocessors – could disable interrupts
 - Currently running code would execute without preemption
 - Generally too inefficient on multiprocessor systems
 - ▶ Operating systems using this not broadly scalable
- Modern machines provide special atomic hardware instructions
 - ▶ Atomic = non-interruptible
 - Either test memory word and set value
 - Or swap contents of two memory words

TestAndSet Instruction

➤ Definition:

```
boolean TestAndSet (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

➤ Solution: Shared Boolean variable *lock*, initialized to false.

```
while (true) {
    while ( TestAndSet (&lock ))
        ; /* do nothing
           // critical section
    lock = FALSE;
        // remainder section
}
```

Swap Instruction

➤ Definition:


```
void Swap (boolean *a, boolean *b)
{
    boolean temp = *a;
    *a = *b;
    *b = temp;
}
```

➤ Solutions: Shared Boolean variable *lock* initialized to FALSE; Each process has a local Boolean variable *key*.

```
while (true) {
    key = TRUE;
    while ( key == TRUE)
        Swap (&lock, &key );
    // critical section
    lock = FALSE;
    // remainder section
}
```

Semaphore

- Synchronization tool that does not require busy waiting
 - integer variable
 - Two standard operations
 - ▶ `S.wait()` \Rightarrow `P()`
 - ▶ `S.signal()` \Rightarrow `V()`
 - Less complicated
- Can only be accessed via two indivisible (atomic) operations



```
wait (S) {  
    while S <= 0; // no-op  
    S--;  
}
```

```
signal (S) {  
    S++;  
}
```

Semaphore as General Synchronization Tool

- **Binary** semaphore – integer value can range only between 0 and 1; can be simpler to implement
 - Also known as **mutex locks**
- **Counting** semaphore – integer value can range over an unrestricted domain
- Can implement a counting semaphore **S** as a binary semaphore
- Provides mutual exclusion
 - Semaphore **S**; // initialized to 1
 - wait (S);
Critical Section
signal (S);

Semaphore Implementation

- Must guarantee that no two processes can execute `wait ()` and `signal ()` on the same semaphore at the same time
- Thus, implementation becomes the critical section problem where the wait and signal code are placed in the critical section.
 - Could have busy waiting in critical section implementation
 - ▶ But implementation code is short
 - ▶ Little busy waiting if critical section rarely occupied
- Note that applications may spend lots of time in critical sections and therefore this is not a good solution.

No Busy waiting Implementation

- With each semaphore there is an associated waiting queue.
 - value (of type integer)
 - a list of processes list (a list of PCBs)

- Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue.
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue.

No Busy waiting Implementation (Cont.)

➤ Implementation of wait:

```
wait (S){  
    value--;  
    if (value < 0) {  
        add this process to waiting queue  
        block();  
    }  
}
```

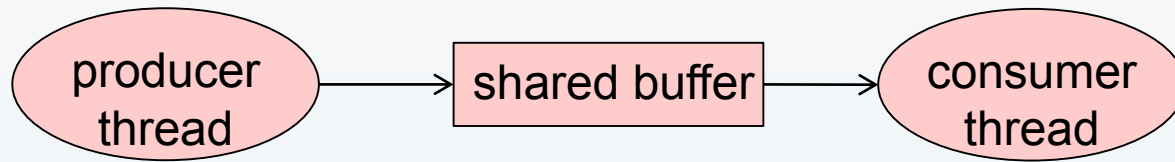
➤ Implementation of signal:

```
Signal (S){  
    value++;  
    if (value <= 0) {  
        remove a process P from the waiting queue  
        wakeup(P);  
    }  
}
```


Classical Problems of Synchronization

- Bounded-Buffer Problem
- Readers and Writers Problem

Bounded-Buffer Problem



➤ Common synchronization pattern:

- Producer waits for slot, inserts item in buffer, and signals consumer
- Consumer waits for item, removes it from buffer, and signals producer

➤ Examples

- **Multimedia processing:**
 - ▶ Producer creates MPEG video frames, consumer renders the frames
- **Graphical user interfaces**
 - ▶ Producer detects mouse clicks, mouse movements, and keyboard hits and inserts corresponding events in buffer.
 - ▶ Consumer retrieves events from buffer and paints the display.

Solving a Bounded-Buffer Problem

```
/* buffer.c - producer-consumer
on 1-element buffer */
#include <pthread.h>
#define N 5
void *producer(void *arg);
void *consumer(void *arg);
struct {
    int buf; /* shared var */
    sem_t full; /* sems */
    sem_t empty; /* sems */
} shared;
```

```
int main() {
    pthread_t tid_producer;
    pthread_t tid_consumer;
    /* initialize the semaphores */
    Sem_init(&shared.empty, 0, 1);
    Sem_init(&shared.full, 0, 0);
    /* create threads and wait */
    pthread_create(&tid_producer, NULL,
                  producer, NULL);
    pthread_create(&tid_consumer, NULL,
                  consumer, NULL);
    pthread_join(tid_producer, NULL);
    pthread_join(tid_consumer, NULL);
    exit(0);
};
```

Solving a Bounded-Buffer Problem

➤ Initially empty=1, full=0.

```
/* producer thread */
void *producer(void *arg) {
    int i, item;
    for (i=0; i<N; i++) { /* produce item */
        item = i;
        printf("produced %d\n", item);
        /* write item to buf */
        P(&shared.empty);
        shared.buf = item;
        V(&shared.full);
    }
    return NULL;
}
```

```
/* consumer thread */
void *consumer(void *arg) {
    int i, item;
    for (i=0; i<N; i++) { /* read item from buf */
        P(&shared.full);
        item = shared.buf;
        V(&shared.empty);
        /* consume item */
        printf("consumed %d\n", item);
    }
    return NULL;
}
```

Readers-Writers Problem

- A data set is shared among a number of concurrent processes
 - Readers – only read the data set
 - Writers – can both read and write.
- Problems – allow multiple readers to read at the same time. Only one single writer can access the shared data at the same time.
 - First readers-writers problem: no reader should wait for other readers to finish simply because a writer is waiting.
 - Second readers-writers problem: once a writer is ready, no new readers may start reading.

First Readers-Writers Problem Solution

- Data set
- Semaphore **mutex** initialized to 1
 - ▶ mutual exclusion for variable *readcount*.
- Semaphore **wrt** initialized to 1
 - ▶ for both reader and writer processes.
- Integer **readcount** initialized to 0.

```
while (true) {  
    wait (wrt) ;  
    // writing is performed  
    signal (wrt) ;  
}
```

Writer process

```
while (true) {  
    wait (mutex) ;  
    readcount ++ ;  
    if (readcount == 1)  
        wait (wrt) ;  
    signal (mutex)  
    // reading is performed  
    wait (mutex) ;  
    readcount -- ;  
    if (readcount == 0)  
        signal (wrt) ;  
    signal (mutex) ;  
}
```

Reader Process

Problems with Semaphores

- Incorrect use of semaphore operations:
 - signal (mutex) wait (mutex)
 - wait (mutex) ... wait (mutex)
 - Omitting of wait (mutex) or signal (mutex) (or both)