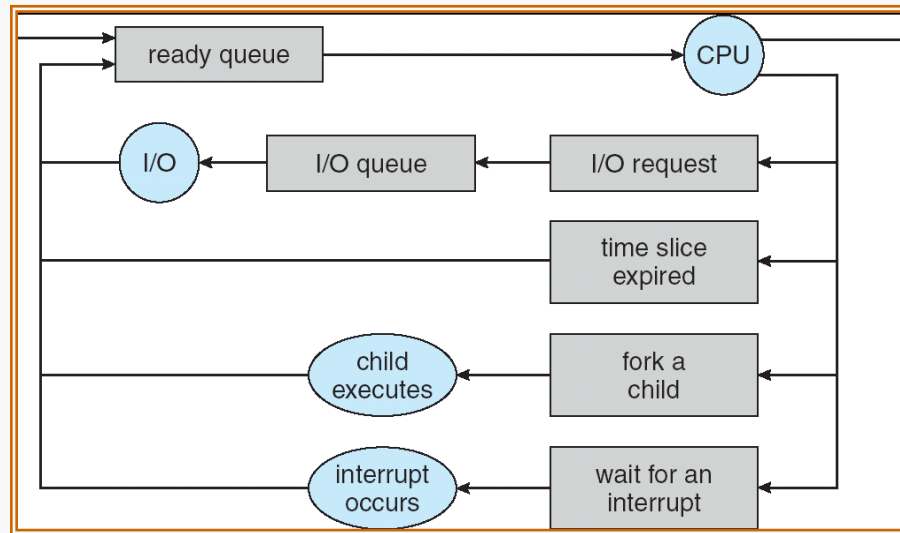


# **CSI 4500 Operating Systems**


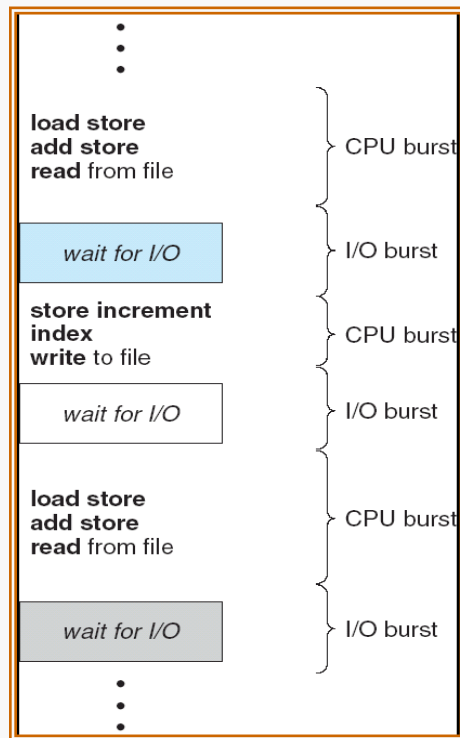
## **CPU Scheduling**

# CPU Scheduling



- life-cycle of a thread shows above
  - Active threads work their way from **Ready** queue to **Running** with possible various **Waiting** queues.
- Question: How is the OS to decide which of several tasks to run next?
  - Obvious queue to worry about is ready queue
- **Scheduling**: deciding which threads are given access to resources from moment to moment.

# Assumption: CPU-I/O Bursts Cycle



- Execution model: programs alternate between bursts of **CPU** and **I/O**
  - Program typically uses the CPU for some period of time, then does I/O, then uses CPU again
  - Each scheduling decision is about which job to give to the CPU for use by its next CPU burst
  - With timeslicing, thread may be forced to give up CPU before finishing current CPU burst

# CPU Scheduler

- Selects from among the processes in memory that are ready to execute, and allocates the CPU to one of them
- CPU scheduling decisions may take place when a process:
  1. Switches from running to waiting state
  2. Terminates
  3. Switches from waiting to ready
  4. Switches from running to ready state

Cooperative scheduling

**Preemptive scheduling**

# Scheduling Criteria

- **CPU utilization** – keep the CPU as busy as possible
- **Throughput** – number of processes that complete their execution per time unit
- **Turnaround time** – amount of time to execute a particular process, from submission till the time of completion.
- **Waiting time** – amount of time a process has been waiting in the ready queue
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced, **not** output (for time-sharing environment)

# Scheduling Policy Goals

- Minimize Response Time
  - Response time is what the user sees:
    - ▶ Time to echo a keystroke in editor
    - ▶ Time to compile a program
    - ▶ Real-time Tasks: Must meet deadlines imposed by World
- Maximize Throughput
  - Throughput related to response time, but not identical:
    - ▶ Minimizing response time will lead to more context switching than if you only maximized throughput
  - Two parts to maximizing throughput
    - ▶ Minimize overhead (for example, context-switching)
    - ▶ Efficient use of resources (CPU, disk, memory, etc)
- Fairness
  - Share CPU among users in some equitable way
  - Fairness is not minimizing average response time:
    - ▶ Better *average* response time by making system *less* fair

# First-Come, First-Served (FCFS) Scheduling

- First-Come, First-Served (FCFS)
  - Also “First In, First Out” (FIFO) or “Run until done”
    - ▶ In early systems, FCFS meant one program scheduled until done (including I/O)
    - ▶ Now, means keep CPU until thread blocks



- Example: 

Process	Burst Time
$P_1$	24
$P_2$	3
$P_3$	3

$P_1$	24
$P_2$	3
$P_3$	3

- Suppose processes arrive in the order:  $P_1, P_2, P_3$   
The Gantt Chart for the schedule is:



- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- Average waiting time:  $(0 + 24 + 27)/3 = 17$
- Average Completion time:  $(24 + 27 + 30)/3 = 27$
- *Convoy effect*: short process behind long process

# Round Robin (RR)




- FCFS Scheme: Potentially bad for short jobs!
  - Depends on submit order
  - If you are first in line at supermarket with milk, you don't care who is behind you, on the other hand...
- Round Robin Scheme
  - Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds
  - After quantum expires, the process is preempted and added to the end of the ready queue.
  - $n$  processes in ready queue and time quantum is  $q \Rightarrow$ 
    - ▶ Each process gets  $1/n$  of the CPU time
    - ▶ In chunks of at most  $q$  time units
    - ▶ No process waits more than  $(n-1)q$  time units
- Performance
  - $q$  large  $\Rightarrow$  FCFS
  - $q$  small  $\Rightarrow$  Interleaved (really small  $\Rightarrow$ ?)
  - $q$  must be large with respect to context switch, otherwise overhead is too high (all overhead)



# Round-Robin Discussion

- Thus, Round-Robin Pros and Cons:
  - Better for short jobs, Fair (+)
  - Context-switching time adds up for long jobs (-)
- How do you choose time slice?
  - What if too big?
    - ▶ Response time suffers
  - What if infinite ( $\infty$ )?
    - ▶ Get back FIFO
  - What if time slice too small?
    - ▶ Throughput suffers!
- Actual choices of timeslice:
  - Initially, UNIX timeslice one second:
    - ▶ Worked ok when UNIX was used by one or two people.
    - ▶ What if three compilations going on? 3 seconds to echo each keystroke!
  - In practice, need to balance short-job performance and long-job throughput:
    - ▶ Typical time slice today is between 10ms – 100ms
    - ▶ Typical context-switching overhead is 0.1ms – 1ms
    - ▶ Roughly 1% overhead due to context-switching



# Comparisons between FCFS and RR

❑ Assuming zero-cost context-switching time, is RR always better than FCFS?

❑ Simple example:

10 jobs, each take 100s of CPU time  
RR scheduler quantum of 1s  
All jobs start at the same time

❑ Completion Times:

Job #	FCFS	RR
1	100	991
2	200	992
...	...	...
9	900	999
10	1000	1000

➤ Both RR and FCFS finish at the same time

➤ Average response time

▶ RR: 4.5

▶ FCFS: 450

Average waiting time

RR: 891

FCFS: 450

❑ Also: Cache state must be shared between all jobs with RR but can be devoted to each job with FIFO

➤ Total time for RR longer even for zero-cost switch!


# What if we Knew the Future?

- Shortest Job First (SJF):
  - Run whatever job has the least amount of computation to do
  - Sometimes called “Shortest Time to Completion First” (STCF)
- Shortest Remaining Time First (SRTF):
  - Preemptive version of SJF: if job arrives and has a shorter time to completion than the remaining time on the current job, immediately preempt CPU
- These can be applied either to a whole program or the current CPU burst of each program
  - Idea is to get short jobs out of the system as soon as possible
  - Big effect on short jobs, only small effect on long ones
  - Result is better average response time
- SJF is optimal – gives minimum average waiting time for a given set of processes

# Example of Non-Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

- SJF (non-preemptive)




- Average waiting time =  $(0 + 6 + 3 + 7)/4 = 4$

# Example of Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

- SJF (preemptive)




- Average waiting time =  $(9 + 1 + 0 + 2)/4 = 3$

# Discussion

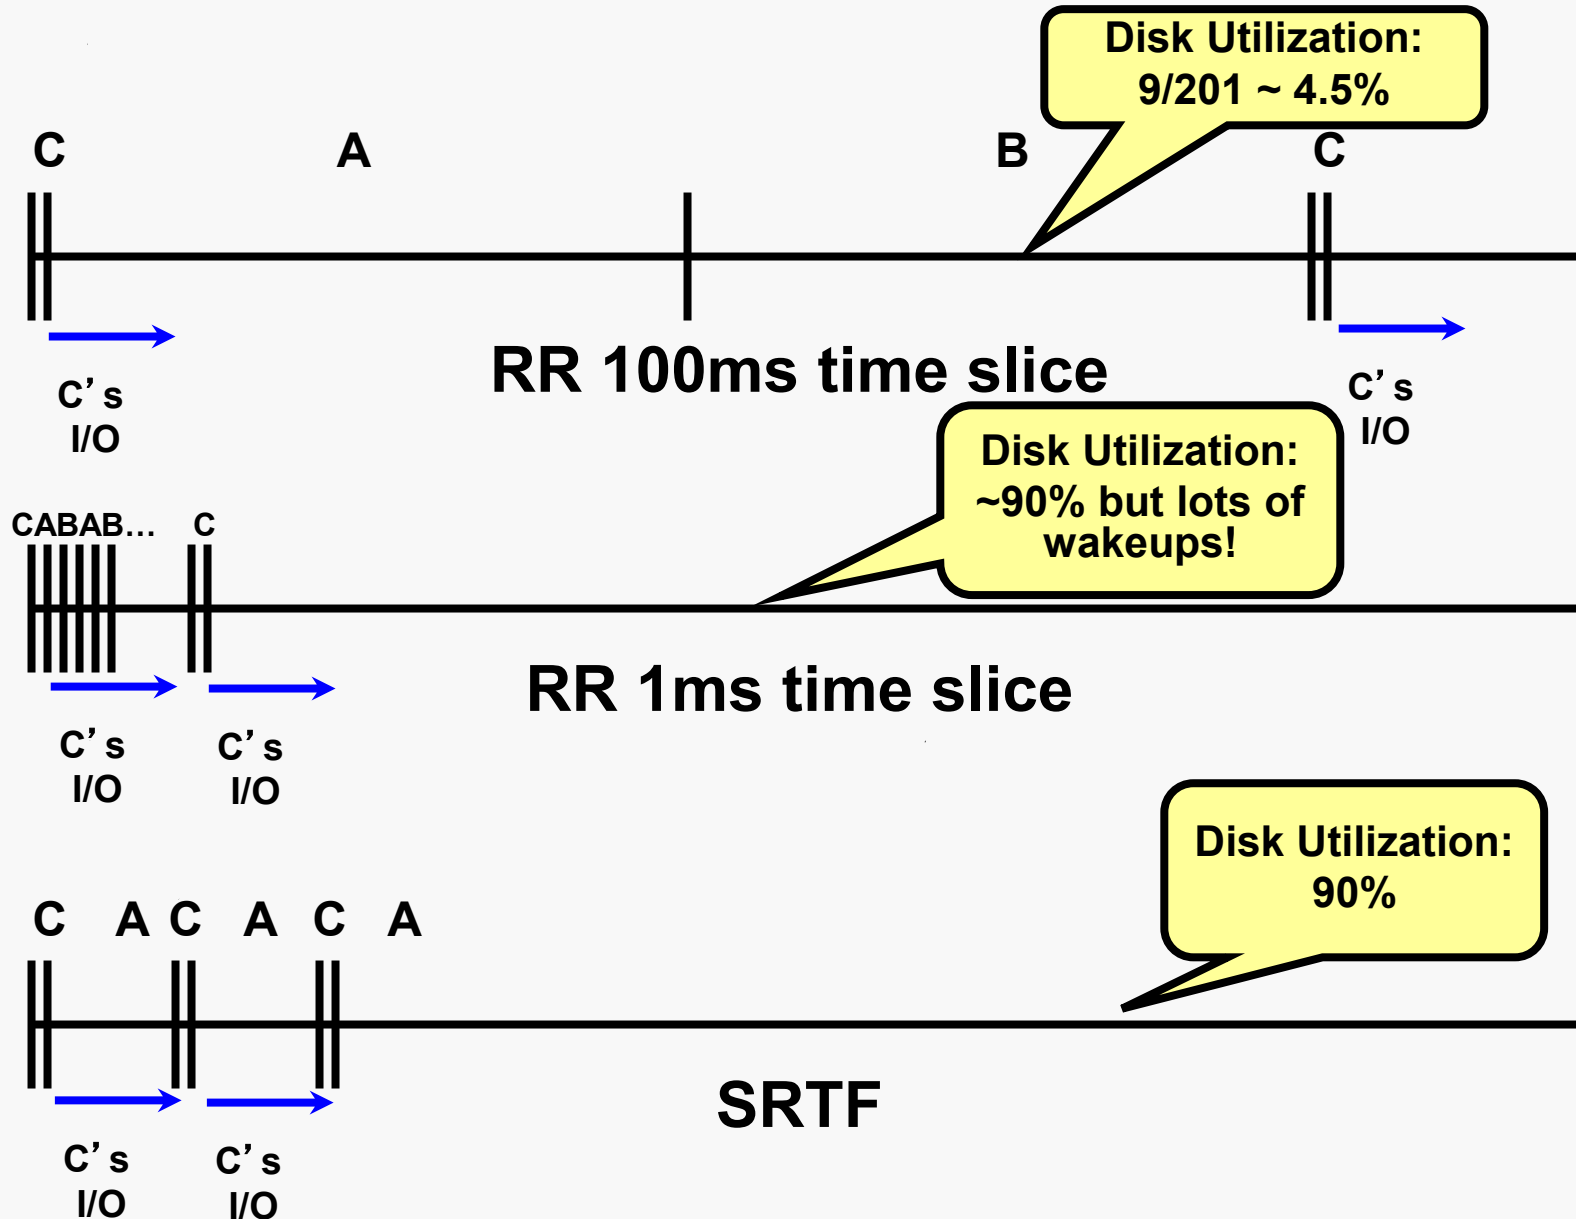
- SJF/SRTF are the best you can do at minimizing average response time
  - Provably optimal (SJF among non-preemptive, SRTF among preemptive)
  - Since SRTF is always at least as good as SJF, focus on SRTF
- Comparison of SRTF with FCFS and RR
  - What if all jobs the same length?
    - ▶ SRTF becomes the same as FCFS
    - ▶ What if jobs have varying length?
    - ▶ SRTF (and RR): short jobs not stuck behind long ones

# Example to illustrate benefits of SRTF



- Three jobs:
  - A,B: both CPU bound, run for weeks
  - C: I/O bound, loop 1ms CPU, 9ms disk I/O
  - If only one at a time, C uses 90% of the disk, A or B could use 100% of the CPU
- With FIFO:
  - Once A or B get in, keep CPU for weeks
- What about RR or SRTF?
  - Easier to see with a timeline

# SRTF Example continued:





# SRTF Further discussion

- Starvation
  - SRTF can lead to starvation if many small jobs!
  - Large jobs never get to run
- Somehow need to predict future
  - How can we do this?
  - Some systems ask the user
    - ▶ When you submit a job, have to say how long it will take
    - ▶ To stop cheating, system kills job if takes too long
  - But: Even non-malicious users have trouble predicting runtime of their jobs
- Bottom line, can't really know how long job will take
  - However, can use SRTF as a yardstick for measuring other policies
  - Optimal, so can't do any better
- SRTF Pros & Cons
  - Optimal (average response time) (+)
  - Hard to predict future (-)
  - Unfair (-)

# Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer  $\equiv$  highest priority)
  - Preemptive
  - nonpreemptive
- SJF is a priority scheduling where priority is the predicted next CPU burst time
- Problem  $\equiv$  **Starvation** – low priority processes may never execute
- Solution  $\equiv$  **Aging** – as time progresses increase the priority of the process [giving more CPU time gradually]

# Challenges

- No universal scheduler
  - Unpredictable workloads
  - Unpredictable environment
- Building a one-size-fits-all scheduler that works well for all is **challenging**!
- Real scheduling algorithms are often more complex than the simple scheduling algorithms we've seen
  - FCFS
  - RR
  - SJF/SRTF

# A solution

- **Multilevel Queue Scheduling:** ready queue is partitioned into multiple queues
  - Multiple queues, each with different priority
    - ▶ Higher priority queues often considered “foreground” tasks
  - Each queue has its own scheduling algorithm
    - ▶ e.g. foreground – RR, background – FCFS
    - ▶ Sometimes multiple RR priorities with quantum increasing exponentially (highest:1ms, next:2ms, next: 4ms, etc)
- Must choose scheduling algorithm to schedule between queues.
  - Fixed priority scheduling for each queue
    - ▶ serve all from highest priority, then next priority, etc.
  - **RR** between queues: Time slice:
    - ▶ each queue gets a certain amount of CPU time
    - ▶ e.g., 70% to highest, 20% next, 10% lowest

# Multilevel Feedback Queue

- A process can move between the various queues
  - aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
  - number of queues
  - scheduling algorithms for each queue
  - method used to determine when to upgrade a process
  - method used to determine when to demote a process
  - method used to determine which queue a process will enter when that process needs service


# Example of Multilevel Feedback Queue

- Three queues:

- $Q_0$  – RR with time quantum 8 milliseconds
- $Q_1$  – RR time quantum 16 milliseconds
- $Q_2$  – FCFS

- Scheduling

- A new job enters queue  $Q_0$  which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue  $Q_1$ .
- At  $Q_1$  job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue  $Q_2$ .



# What about Fairness?

- What about fairness?
  - Strict fixed-priority scheduling between queues is unfair (run highest, then next, etc):
    - ▶ long running jobs may never get CPU
    - ▶ In Multics, shut down machine, found 10-year-old job
  - Must give long-running jobs a fraction of the CPU even when there are shorter jobs to run
  - Tradeoff: fairness gained by hurting average response time!

# What about Fairness?

- How to implement fairness?
  - Could give each queue some fraction of the CPU
    - ▶ What if one long-running job and 100 short-running ones?
    - ▶ Like express lanes in a supermarket—sometimes express lanes get so long, get better service by going into one of the other lines
  - Could increase priority of jobs that don't get service
    - ▶ What is done in UNIX
    - ▶ This is ad hoc—what rate should you increase priorities?
    - ▶ And, as system gets overloaded, no job gets CPU time, so everyone increases in priority  $\Rightarrow$  Interactive jobs suffer



# Linux Scheduling

- Avoid starvation
- Boost interactivity
  - **Fast response** to user despite high load
  - Achieved by inferring interactive processes and dynamically increasing their priorities
- SMP goals
  - **Scale** well w.r.t number of processes
    - ▶  $O(1)$  scheduling overhead
  - **Load balance**: no CPU should be idle if there is work
  - **CPU affinity**: no random bouncing of processes

# Linux Scheduling Algorithm Overview

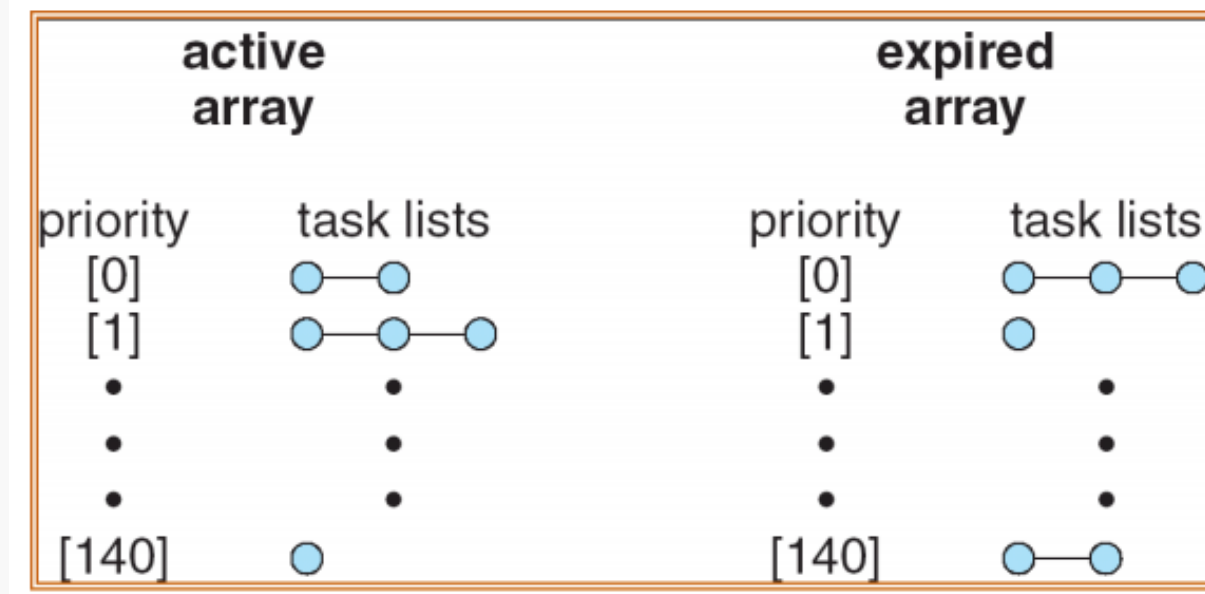
- Multilevel queue scheduler
  - Each queue associated with a **priority**
  - A process's priority may be adjusted **dynamically**
- Two classes of processes
  - Real-time processes: always schedule highest priority processes
    - ▶ FCFS (**SCHED\_FIFO**) or RR (**SCHED\_RR**) for processes with same priority
  - Normal processes: priority with aging
    - ▶ RR for processes with same priority (**SCHED\_NORMAL**)
    - ▶ Aging is implemented efficiently

# Priority partition

- Total 140 priorities [0, 140)
  - Smaller integer == higher priority
  - Real-time: [0, 100)
  - Normal: [100, 140)
- **MAX\_PRIO** and **MAX\_RT\_PRIO**
  - `include/linux/sched.h`

# Runqueue data structure

- kernel/sched.c
- struct prio\_array
  - Array of priority queues
- struct runqueue
  - Two arrays, **active** and **expired**



# Scheduling Algorithm

1. Find highest priority non-empty queue in `rq->active`; if none, simulate aging by swapping `active` and `expired`
2. `next` = first process on that queue
3. Adjust `next`'s priority
  1. Dynamically increase priority of interactive process
4. Context switch to `next`
5. When `next` used up its time slice, insert `next` to the right queue and call `schedule` again

`schedule()` in `kernel/sched.c`

# Summary

- Scheduling problem
  - Given a set of processes that are ready to run
  - Which one to select *next*
- Scheduling criteria
  - CPU utilization, Throughput, Turnaround, Waiting, Response
  - Predictability: variance in any of these measures
- Scheduling algorithms
  - FCFS, SJF, SRTF, RR
  - Multilevel (Feedback-)Queue Scheduling
- The best schemes are adaptive.