

# **CSI 4500**

## **Operating Systems**

### **Deadlock**

# Goals for Today

- Understand Deadlock
- Understand Resource Allocation Graph
- Study techniques for
  - Deadlock Prevention
  - Deadlock Avoidance
  - Deadlock Detection


# Definitions

## ■ Starvation: thread waits indefinitely

- ▶ Example, low-priority thread waiting for resources constantly in use by high-priority threads

## ■ Deadlock: circular waiting for resources

- ▶ Thread A owns Res 1 and is waiting for Res 2  
Thread B owns Res 2 and is waiting for Res 1



## ■ Deadlock $\Rightarrow$ Starvation but not vice versa

- ▶ Starvation can end (but doesn't have to)
- ▶ Deadlock can't end without external intervention

# Conditions for Deadlock

<u>Thread A</u>	<u>Thread B</u>
<code>semX.P ();</code>	<code>semY.P ();</code>
<code>semY.P ();</code>	<code>semX.P ();</code>
<code>semY.V ();</code>	<code>semX.V ();</code>
<code>semX.V ();</code>	<code>semY.V ();</code>

- Deadlock won't always happen with this code
  - Have to have exactly the right timing order
- Deadlocks occur with multiple resources
  - Can't solve deadlock for each resource independently
- Example: System with 1 disk drive, 1 printer and two threads
  - Each thread needs both of them to function
  - Thread 1 gets disk and Thread 2 gets the printer

# Four requirements for Deadlock

## ■ Mutual exclusion

- Only one thread at a time can use a resource.

## ■ Hold and wait

- Thread holding at least one resource is waiting to acquire additional resources held by other threads

## ■ No preemption

- Resources are released only voluntarily by the thread holding the resource, after thread is finished with it

## ■ Circular wait

- There exists a set  $\{T_1, \dots, T_n\}$  of waiting threads
  - ▶  $T_1$  is waiting for a resource that is held by  $T_2$
  - ▶  $T_2$  is waiting for a resource that is held by  $T_3$
  - ▶ ...
  - ▶  $T_n$  is waiting for a resource that is held by  $T_1$

# System Model & Resource-Allocation Graph

## ■ System Model

- A set of processes  $P_1, P_2, \dots, P_n$
- Resource types  $R_1, R_2, \dots, R_m$ 
  - E.g., CPU cycles, memory space, I/O devices
- Each resource type  $R_i$  has  $W_i$  instances.
- Each process utilizes a resource as follows:
  - ▶ Request ()  $\rightarrow$  Use ()
  - $\rightarrow$  Release ()

## ■ Resource-Allocation Graph:

- V is partitioned into two types:

▶  $P = \{P_1, P_2, \dots, P_n\}, R = \{R_1, R_2, \dots, R_m\}$


- E is partitioned into two types:

▶ request edge – directed edge

–  $P_1 \rightarrow R_j$

▶ assignment edge – directed edge  $R_j \rightarrow P_i$

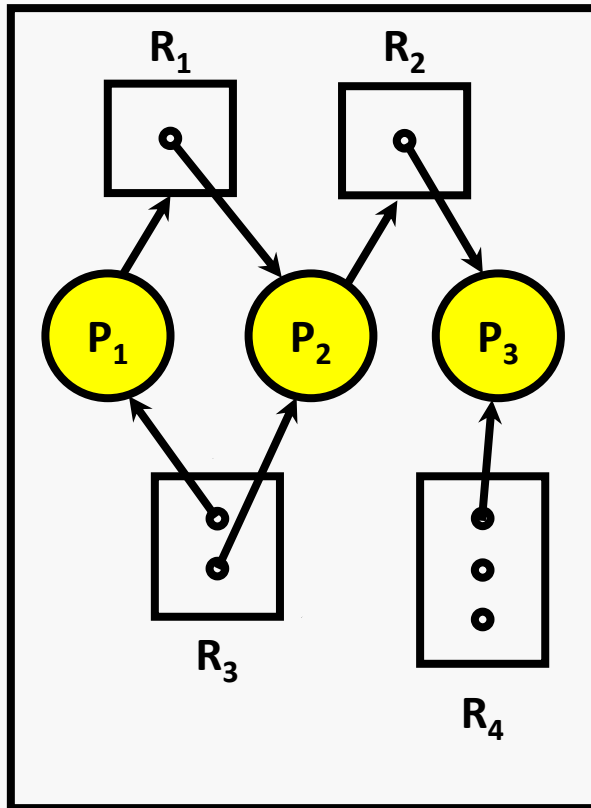
### Symbols



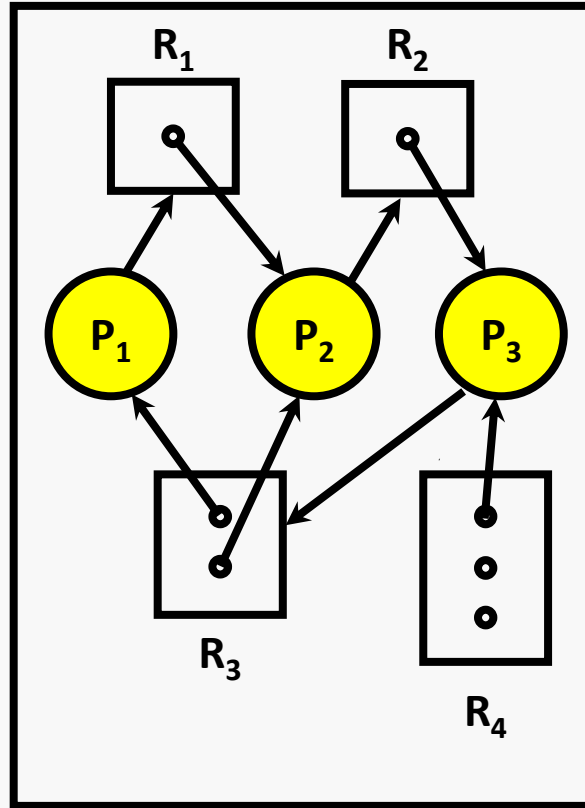
# Resource Allocation Graph Examples

## ■ Recall:


- **request** edge – directed edge  $P_1 \rightarrow R_j$
- **assignment** edge – directed edge  $R_j \rightarrow T_i$



Simple Resource  
Allocation Graph



Allocation Graph  
With Deadlock



Allocation Graph  
With Cycle, but  
No Deadlock

# Initial Found

- If graph contains no cycles  $\Rightarrow$  no deadlock.
- If graph contains a cycle  $\Rightarrow$ 
  - if only one instance per resource type, then deadlock.
  - if several instances per resource type, possibility of deadlock.



# Methods for Handling Deadlocks

- Ensure that system will **NEVER** enter a deadlock
  - Need to monitor all lock acquisitions
  - Selectively deny those that **MIGHT** lead to deadlock
- **ALLOW** system to enter deadlock and then recover
  - Requires deadlock detection algorithm
  - Some technique for forcibly preempting resources and/or terminating tasks
- **Ignore** the problem and pretend that deadlocks never occur in the system
  - Used by most operating systems, including UNIX, windows

# Deadlock Prevention

## ■ Mutual Exclusion

- not required for sharable resources; must hold for nonsharable resources.

## ■ Hold and Wait – must guarantee that whenever a process requests a resource, it does not hold any other resources.

- Require process to request and be allocated all its resources before it begins execution
- Allow process to request resources only when the process has none.
- Low resource utilization; starvation possible.

## ■ No Preemption

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released.
- Preempted resources are added to the list of resources for which the process is waiting.
- Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting.

## ■ Circular Wait

- a total ordering of all resource types
- require that each process requests resources in an ordering strategy

# Deadlock Avoidance

## ■ Resource-allocation state

- the number of available and allocated resources
- the maximum demands of the processes.


- The deadlock-avoidance algorithm dynamically examines the **resource-allocation state** to ensure that there can never be a circular-wait condition.

# Safe State

- System is in **safe state** if there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of ALL the processes in the systems such that for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by the currently available resources + resources held by all the  $P_j$ , with  $j < i$ .
- That is:
  - If  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished.
  - When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate.
  - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on.

# Basic Facts

- If a system is in safe state  $\Rightarrow$  no deadlocks.
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock.
- Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state.




# Deadlock Avoidance algorithms

- Single instance of a resource type.
  - Use a resource-allocation graph
- Multiple instances of a resource type.
  - Use the banker's algorithm

# Resource-Allocation Graph Scheme


- **Claim edge**  $P_i \rightarrow R_j$  indicated that process  $P_j$  may request resource  $R_j$  in the future.
  - represented by a dashed line.
- **Claim edge** converts to **request edge** when a process requests a resource.
- **Request edge** converted to an **assignment edge** when the resource is allocated to the process.
- When a resource is released by a process, **assignment edge** reconverts to a **claim edge**.
- Resources must be claimed *a priori* in the system.



# Resource-Allocation Graph Algorithm

- Suppose that process  $P_i$  requests a resource  $R_j$
- The request can be granted only if converting the **request edge** to an **assignment edge** does not result in the formation of a cycle in the resource allocation graph

●  $P_i \rightarrow R_j \rightarrow R_j \rightarrow P_i$



Resource Allocation leading to an Unsafe State  
In Resource-Allocation Graph



# Banker's Algorithm

- For one resource type, there are multiple instances.
- Each process must a priori claim maximum use.
  - This #  $\leq$  total # of resources in the system
- When a process requests a resource it may have to wait.
- When a process gets all its resources it must return them in a finite amount of time.
- When system will determine whether the allocation of these resources will leave the system in a safe state.

# Data Structures for the Banker's Algorithm

Let  $n$  = number of processes, and  $m$  = number of resources types.

- **Available**: Vector of length  $m$ . If  $available[j] = k$ , there are  $k$  instances of resource type  $R_j$  available.
- **Max**:  $n \times m$  matrix. If  $Max[i,j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$ .
- **Allocation**:  $n \times m$  matrix. If  $Allocation[i,j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$ .
- **Need**:  $n \times m$  matrix. If  $Need[i,j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task.

$$Need[i,j] = Max[i,j] - Allocation[i,j].$$



# Safety Algorithm

1. Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , respectively. Initialize:

*Work = Available*

*Finish [i] = false for  $i = 0, 1, \dots, n-1$ .*

2. Find an  $i$  such that both:

(a) *Finish[i] = false*

(b)  $Need_i \leq Work$

If no such  $i$  exists, go to step 4.

3.  $Work = Work + Allocation_i$   
*Finish[i] = true*  
go to step 2.

4. If *Finish [i] == true* for all  $i$ , then the system is in a safe state.

# Resource-Request Algorithm for Process $P_i$

- Definition: *Request<sub>i</sub>* = request vector for process  $P_i$ .
  - ▶ If  $Request_i[j] = k$  then process  $P_i$  wants  $k$  instances of resource type  $R_j$ .
- 1. If  $Request_i \leq Need_i$  go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
- 2. If  $Request_i \leq Available$ , go to step 3. Otherwise  $P_i$  must wait, since resources are not available.
- 3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:
  - $Available = Available - Request_i$
  - $Allocation_i = Allocation_i + Request_i$
  - $Need_i = Need_i - Request_i$
- 4. *Run Safety algorithm*
  - If safe  $\Rightarrow$  the resources are allocated to  $P_i$ .
  - If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored

# Example of Banker's Algorithm

- 5 processes  $P_0$  through  $P_4$ ; 3 resource types: A (10 instances), B (5 instances), and C (7 instances).
- Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 5 3	3 3 2
$P_1$	2 0 0	3 2 2	
$P_2$	3 0 2	9 0 2	
$P_3$	2 1 1	2 2 2	
$P_4$	0 0 2	4 3 3	

	<u>Need</u>
	A B C
$P_0$	7 4 3
$P_1$	1 2 2
$P_2$	6 0 0
$P_3$	0 1 1
$P_4$	4 3 1

- The content of the matrix *Need* is defined as follows.

$$\text{Need} = \text{Max} - \text{Allocation}.$$

- The system is in a safe state
  - Sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfies safety criteria.

# Example: $P_1$ Request (1,0,2)

- Check that Request  $\leq$  Available, True or False?

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	2 3 0
$P_1$	3 0 2	0 2 0	
$P_2$	3 0 2	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

- Executing safety algorithm shows that sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement.
- Can request for (3,3,0) by  $P_4$  be granted?
- Can request for (0,2,0) by  $P_0$  be granted?

# Deadlock Detection

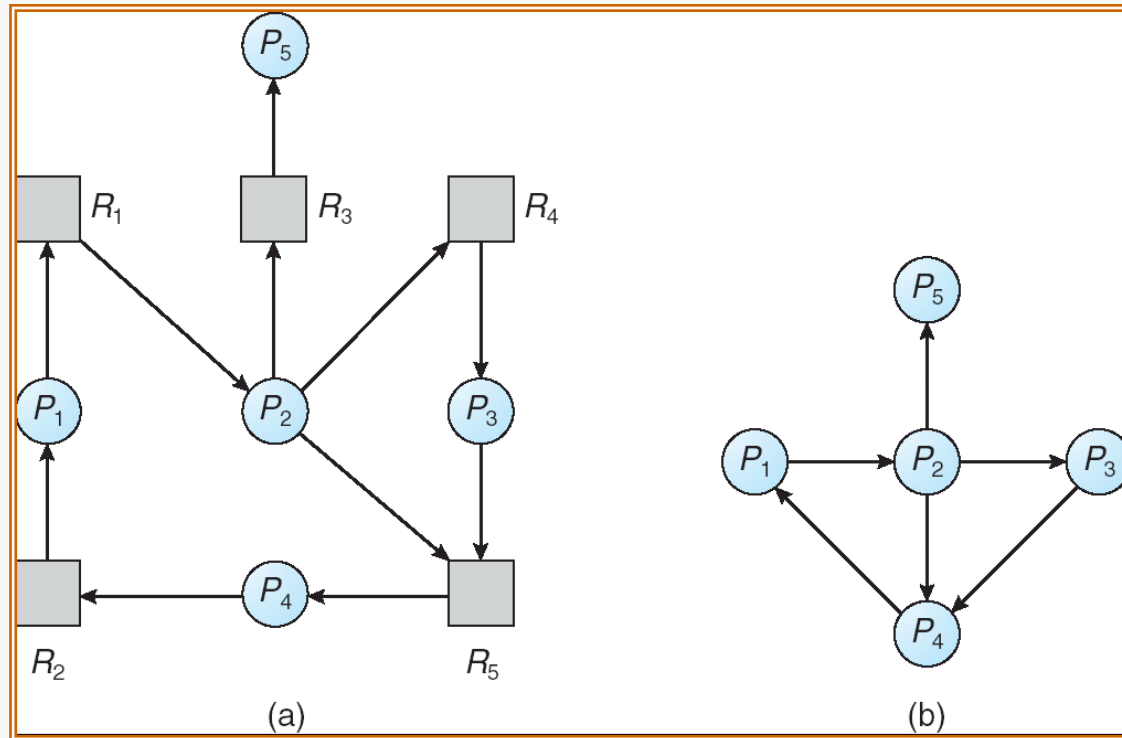
- Allow system to enter deadlock state
- Detection algorithm
  - Single Instance of Resource Type
  - Multiple Instances of Resource Type
- Recovery scheme

# Single Instance of Resource Type

- Maintain *wait-for* graph
  - Nodes are processes.
    - ▶ Remove the resource nodes from resource-allocation graph
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$ .
- Periodically invoke an algorithm that searches for a cycle in the graph.
  - If there is a cycle, there exists a deadlock.
- An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations, where  $n$  is the number of vertices in the graph.



# Resource-Allocation Graph and Wait-for Graph



Resource-Allocation Graph

Corresponding wait-for graph

# Several Instances of a Resource Type

- **Available:** A vector of length  $m$  indicates the number of available resources of each type.
- **Allocation:** An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process.
- **Request:** An  $n \times m$  matrix indicates the current request of each process. If  $Request[i, j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .

# Detection Algorithm

1. Let *Work* and *Finish* be vectors of length  $m$  and  $n$ , respectively Initialize:
  - (a) *Work* = *Available*
  - (b) For  $i = 1, 2, \dots, n$ , if  $Allocation_i \neq 0$ , then  
 $Finish[i] = false$ ; otherwise,  $Finish[i] = true$ .
2. Find an index  $i$  such that both:
  - (a)  $Finish[i] == false$
  - (b)  $Request_i \leq Work$If no such  $i$  exists, go to step 4.
3.  $Work = Work + Allocation_i$   
 $Finish[i] = true$   
go to step 2.
4. If  $Finish[i] == false$ , for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state.  
Moreover, if  $Finish[i] == false$ , then  $P_i$  is deadlocked.

**Algorithm requires an order of  $O(m \times n^2)$  operations to detect whether the system is in deadlocked state.**

# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances).
- Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	0 0 0	0 0 0
$P_1$	2 0 0	2 0 2	
$P_2$	3 0 3	0 0 0	
$P_3$	2 1 1	1 0 0	
$P_4$	0 0 2	0 0 2	

- Sequence  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$  will result in  $Finish[i] = \text{true}$  for all  $i$ .

## Example (Cont.)

- $P_2$  requests an additional instance of type C.

	<u>Request</u>		
	A	B	C
$P_0$	0	0	0
$P_1$	2	0	1
$P_2$	0	0	1
$P_3$	1	0	0
$P_4$	0	0	2

- State of system?

- Can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes' requests.
- Deadlock exists, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$ .

# Detection-Algorithm Usage

- When, and how often, to invoke depends on:
  - How often a deadlock is likely to occur?
  - How many processes will need to be rolled back?
    - ▶ one for each disjoint cycle
- If detection algorithm is invoked arbitrarily, there may be many cycles in the resource graph and so we would not be able to tell which of the many deadlocked processes “caused” the deadlock.

# Recovery from Deadlock: Process Termination

- Abort all deadlocked processes.
- Abort one process at a time until the deadlock cycle is eliminated.
- In which order should we choose to abort?
  - Priority of the process.
  - How long process has computed, and how much longer to completion.
  - Resources the process has used.
  - Resources process needs to complete.
  - How many processes will need to be terminated.
  - Is process interactive or batch?

# Recovery from Deadlock: Resource Preemption

- Selecting a victim – minimize cost.
- Rollback – return to some safe state, restart process for that state.
- Starvation – same process may always be picked as victim, include number of rollback in cost factor.



# Summary

- Starvation vs. Deadlock
  - Starvation: thread waits indefinitely
  - Deadlock: circular waiting for resources
- Four conditions for deadlocks
  - Mutual exclusion
    - ▶ Only one thread at a time can use a resource
  - Hold and wait
    - ▶ Thread holding at least one resource is waiting to acquire additional resources held by other threads
  - No preemption
    - ▶ Resources are released only voluntarily by the threads
  - Circular wait
    - ▶  $\exists$  set  $\{T_1, \dots, T_n\}$  of threads with a cyclic waiting pattern
- Techniques for addressing Deadlock
  - Allow system to enter deadlock and then recover
  - Ensure that system will *never* enter a deadlock
  - Ignore the problem and pretend that deadlocks never occur in the system
- Deadlock detection
  - Attempts to assess whether waiting graph can ever make progress