

CSI 4500 Operating Systems

File System

Today's Objectives

■ File System Motivation

- To explain the function of file systems
- To describe the interfaces to file systems
- To discuss file-system design tradeoffs, including access methods, file sharing, file locking, and directory structures
- To explore file-system protection

■ To describe the details of implementing

- Local file systems
- Directory structures

■ To discuss block allocation and free-block algorithms and trade-offs

Storing Information

■ So far...

- We have discussed processor, memory, I/O

■ How do we make stored information usable?

■ Applications can store information in the process address space

■ Why permanent storage?

- Size is limited to size of virtual address space
 - ▶ May not be sufficient for search engines, banking, etc.
- The data is lost when the application terminates
 - ▶ Even when computer crashes!
- Multiple process might want to access the same data
 - ▶ Imagine a telephone directory part of one process


File Systems

- 3 criteria for long-term information storage:
 - Should be able to store very large amount of information
 - Information must survive the processes using it
 - Should provide concurrent access to multiple processes
- Solution:
 - Store information on disks in units called **files**
 - Files are persistent, and only owner can explicitly delete it
 - Files are managed by the OS
- File Systems: How the OS manages files!

File System

- **File System:** Layer of OS that transforms block interface of disks (or other block devices) into Files, Directories, etc.
- **File System Components**
 - **Disk Management:** collecting disk blocks into files
 - **Naming:** Interface to find files by name, not by blocks
 - **Protection:** Layers to keep data secure
 - **Reliability/Durability:** Keeping of files durable despite crashes, media failures, attacks, etc
- **User vs. System View of a File**
 - User's view: Durable Data Structures
 - System's view (system call interface): Collection of Bytes (UNIX)
 - ▶ Doesn't matter to system what kind of data structures you want to store on disk!
 - System's view (inside OS):
 - ▶ Collection of blocks (a block is a logical transfer unit, while a sector is the physical transfer unit)
 - ▶ Block size \geq sector size; in UNIX, block size is 4KB

Translating from User to System View



■ What happens if user says: give me bytes 2—12?

- Fetch block corresponding to those bytes
- Return just the correct portion of the block

■ What about: write bytes 2—12?

- 1) Fetch block 2) Modify portion 3) Write out Block

■ Everything inside File System is in whole size blocks

- For example, `getc()`, `putc()` \Rightarrow buffers something like 4096 bytes, even if interface is one byte at a time

■ From now on, file is a collection of blocks (i.e. systems view inside OS)

File System Patterns

■ How do users access files?

- Sequential Access

- ▶ bytes read in order (“give me the next X bytes, then give me next, etc”)

- Random Access

- ▶ read/write element out of middle of array (“give me bytes i - j”)

■ What are file sizes?

- Most files are small (for example, .login, .c, .o, .class files, etc)

- Few files are large (for example, core files, etc.)

- Large files use up most of the disk space and bandwidth to/from disk

- ▶ May seem contradictory, but a few enormous files are equivalent to an immense # of small files

File Concept

- A **File** is a named collection of related information that is recorded on secondary storage.
 - Contiguous logical address space

- File represents:
 - Data
 - ▶ Numeric, character, binary
 - ▶ Payroll records, graphic images, sound records, etc.
 - Program
 - ▶ Source programs, object programs, executable programs

File Attributes

■ File Specific information maintained by OS

- File size, modification date, creation time, etc.
- Varies a lot across different OSes
- Information about files are kept in the directory structure, which is maintained on the disk

■ Some Examples:

- **Name** – only information kept in human-readable form
- **Identifier** – unique tag (number) identifies file within file system
- **Type** – needed for systems that support different types
- **Location** – pointer to file location on device
- **Size** – current file size
- **Protection** – controls who can do reading, writing, executing
- **Time, date, and user identification** – data for protection, security, and usage monitoring

File Operations

- **File is an Abstract Data Type**
- **Some basic file operations:**
 - **Create a file**
 - ▶ Find space in FS, add an new entry in the directory
 - **Write a file**
 - ▶ Search file name in the directory to find its location
 - ▶ Information to be written
 - ▶ Write pointer
 - **Read a file**
 - ▶ Search file name in the directory to find its location
 - ▶ Location to put the read results
 - ▶ Read pointer
 - **Reposition within file**
 - ▶ Current-file-position pointer
 - **Delete**
 - ▶ Delete entry from the directory
 - **Truncate**
 - ▶ Keep file attributes except the file length

Open and Close Files

- **Open(F_i)** – search the directory structure on disk for entry F_i , and move the content of entry to memory
 - File pointer: pointer to last read/write location, per process that has the file open
 - File-open count: counter of number of times a file is open
 - Disk location of the file: cache of data access information
 - Access rights: per-process access mode information

- **Close (F_i)** – move the content of entry F_i in memory to directory structure on disk
 - Remove from the open-file table

File Naming


- Motivation: Files abstract information stored on disk
 - You do not need to remember block, sector, ...
 - We have human readable names
- How does it work?
 - Process creates a file, and gives it a name
 - ▶ Other processes can access the file by that name
 - Naming conventions are OS dependent
 - ▶ Usually names as long as 255 characters is allowed
 - ▶ Digits and special characters are sometimes allowed
 - ▶ MS-DOS and Windows are not case sensitive, UNIX family is

File Structure

- Logical record vs physical block
 - Packing a number of logical records into physical block
- Some Example Structures:
 - None - sequence of words, bytes
 - ▶ Unix defines all files to be simply streams of bytes.
 - Simple record structure
 - ▶ Lines, Fixed length, Variable length
 - Complex Structures
 - ▶ Formatted document, Relocatable load file
- Can simulate last two with first method by inserting appropriate control characters

A Typical File-system Organization

- Could use entire disk space for a FS, but
 - A system could have multiple FSes
 - Want to use some disk space for swap space
- Storage Structure
 - A disk is divided into **partitions/slices**.
 - **Directory** records information
 - ▶ Name, location, size, and type



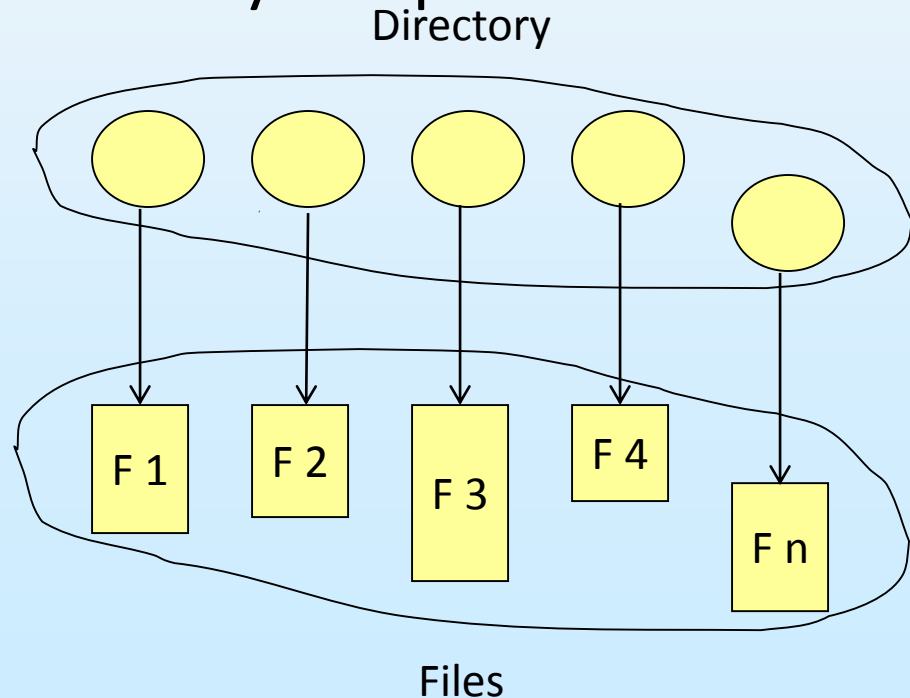
Directory Structure

- A collection of nodes containing information about all files.

- Both the directory structure and the files reside on disk.
- Symbol table

- How to structure the directory to optimize all of the following operations:

- Search a for file
- Create a file
- Delete a file
- List directory
- Rename a file
- Traversing the FS




Organize the Directory (Logically) to Obtain

- Efficiency – locating a file quickly
- Naming – convenient to users
 - Two users can have same name for different files
 - The same file can have several different names
- Grouping – logical grouping of files by properties
 - e.g., all Java programs, all games, ...

Single-Level Directory

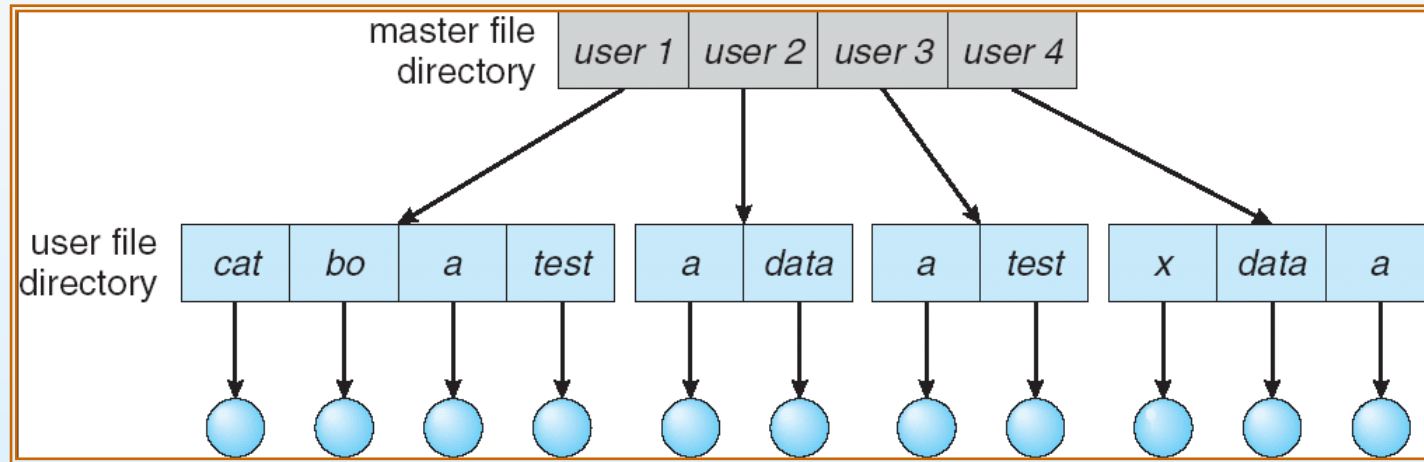
- A single directory for all users
 - Called root directory



- Pros: simplicity, ability to quickly locate files
- Cons: inconvenient naming
 - Naming problem
 - ▶ File name must be unique
 - ▶ remembering all
 - Grouping problem

Two-Level Directory

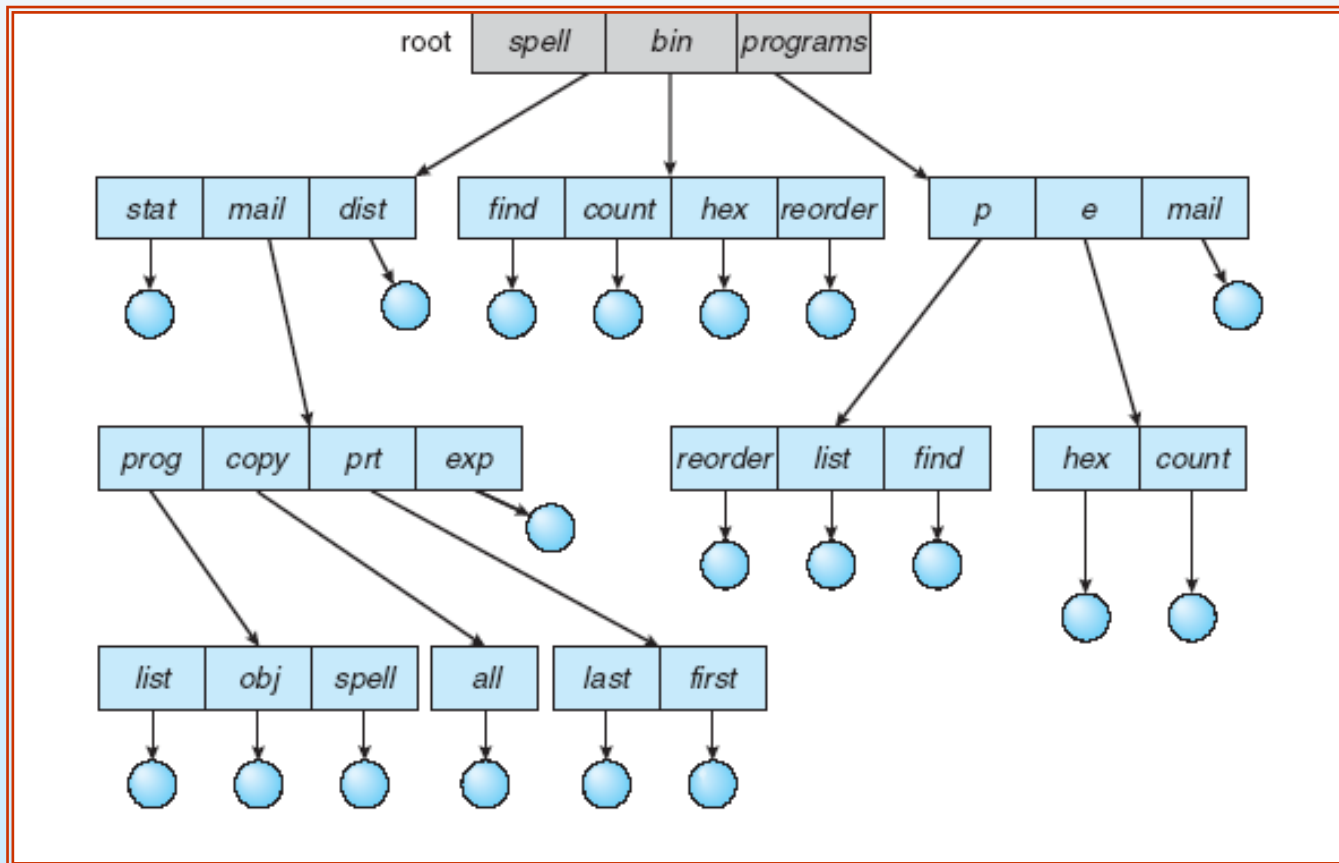
- Separate directory for each user



- Solve name collision, but what if user has lots of files
- Files need to be addressed by path names
 - Allow user's access to other user's files
 - Need for a search path (for example, locating system files)
 - ▶ Efficient searching ?

Tree-Structured Directories

- Directory is now a tree of arbitrary height
 - Directory contains files and subdirectories
 - A bit in directory entry differentiates files from subdirectories



Tree-Structured Directories (Cont)

- **Absolute** or **relative** path name
- Creating a new file is done in current directory
- Delete a file


`rm <file-name>`

- Creating a new subdirectory is done in current directory

`mkdir <dir-name>`

Example: if in current directory `/mail`

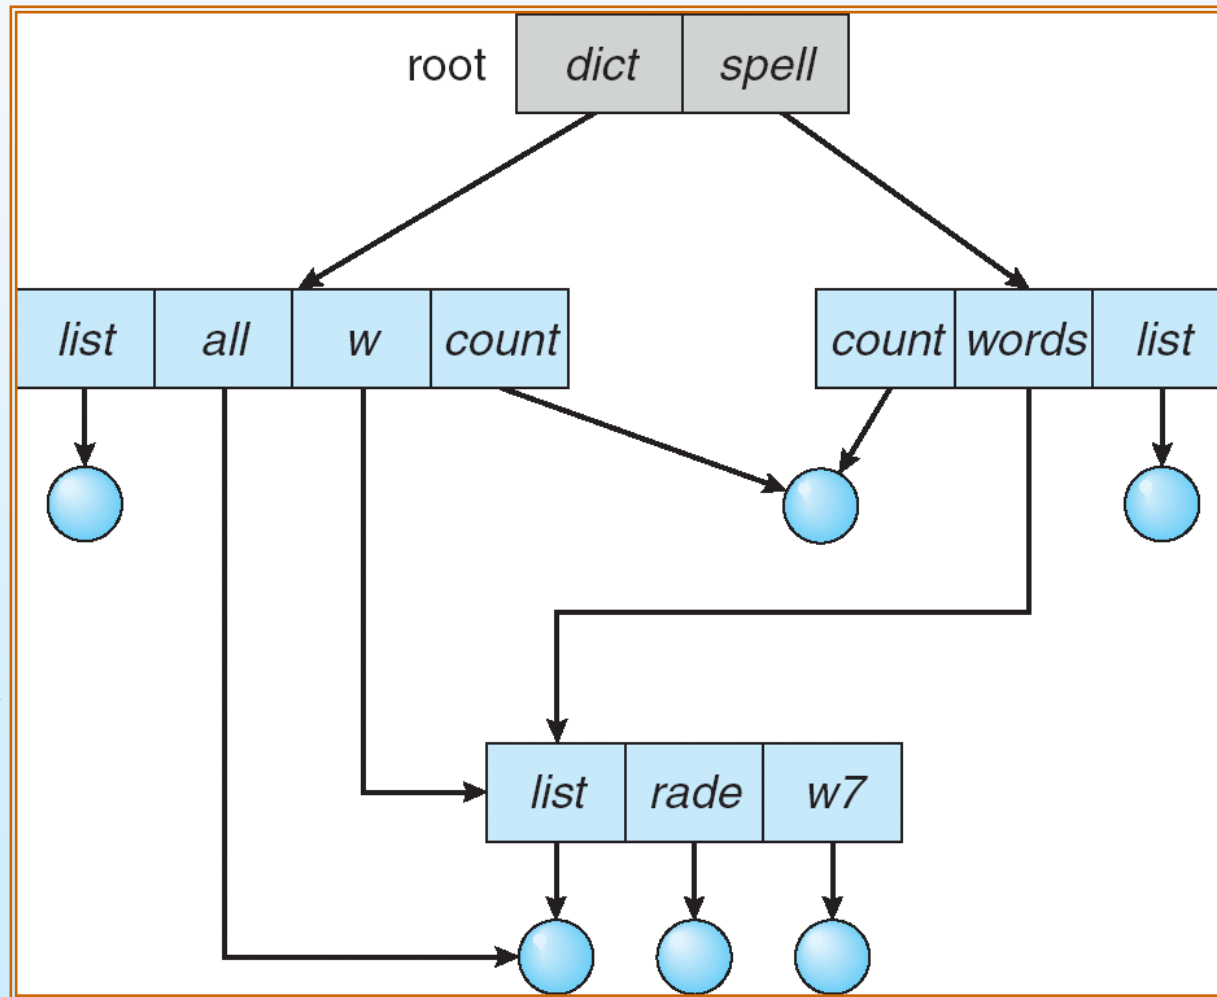
`mkdir count`



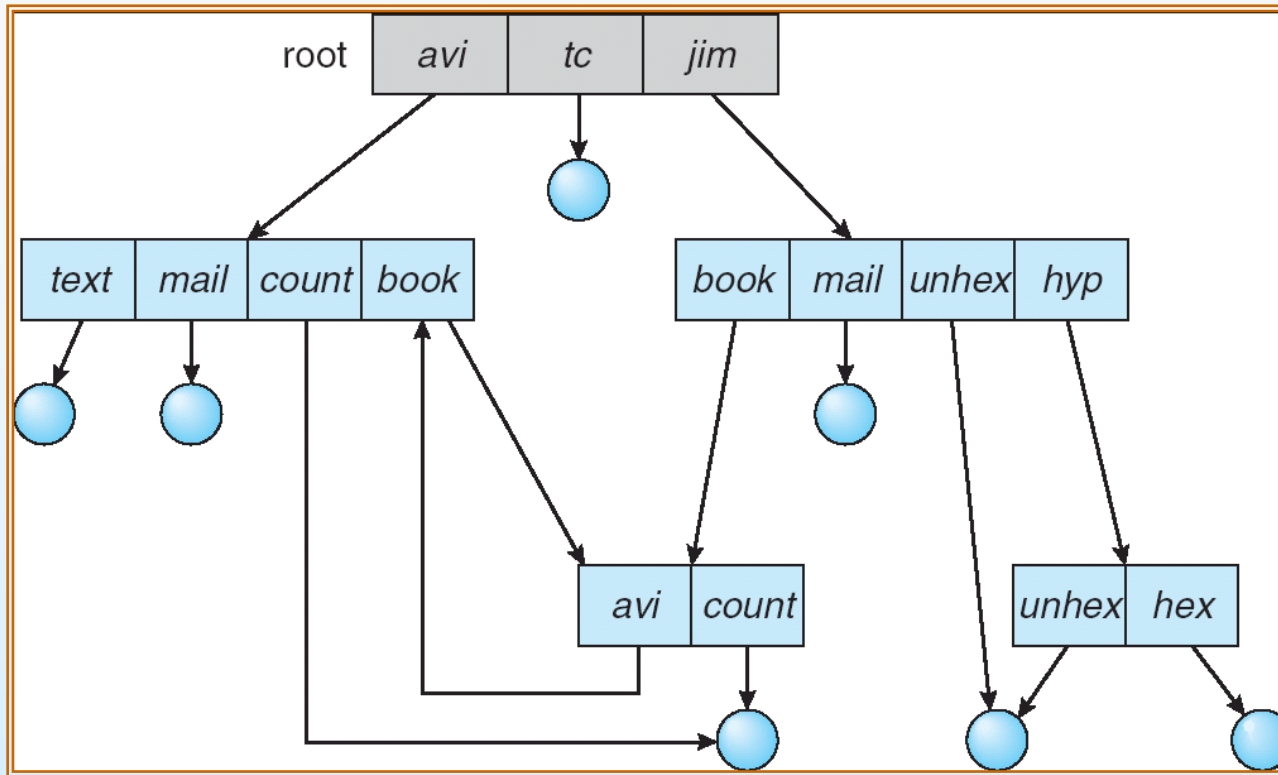
Deleting “mail” \Rightarrow deleting the entire subtree rooted by “mail”

Acyclic-Graph Directories

- Have shared subdirectories and files



General Graph Directory



■ How do we guarantee no cycles?

- Allow only links to file not subdirectories
- Garbage collection
- Every time a new link is added use a cycle detection algorithm to determine whether it is OK

File Sharing


- Sharing of files on multi-user systems is desirable
 - **User IDs** identify users, allowing permissions and protections to be per-user
 - **Group IDs** allow users to be in groups, permitting group access rights
- Sharing may be done through a **protection** scheme
- On distributed systems, files may be shared across a network
- Network File System (NFS) is a common distributed file-sharing method

Access Lists and Groups

- Mode of access: read, write, execute
- Three classes of users

			RWX
a) owner access	7	⇒	1 1 1
			RWX
b) group access	6	⇒	1 1 0
			RWX
c) public access	1	⇒	0 0 1

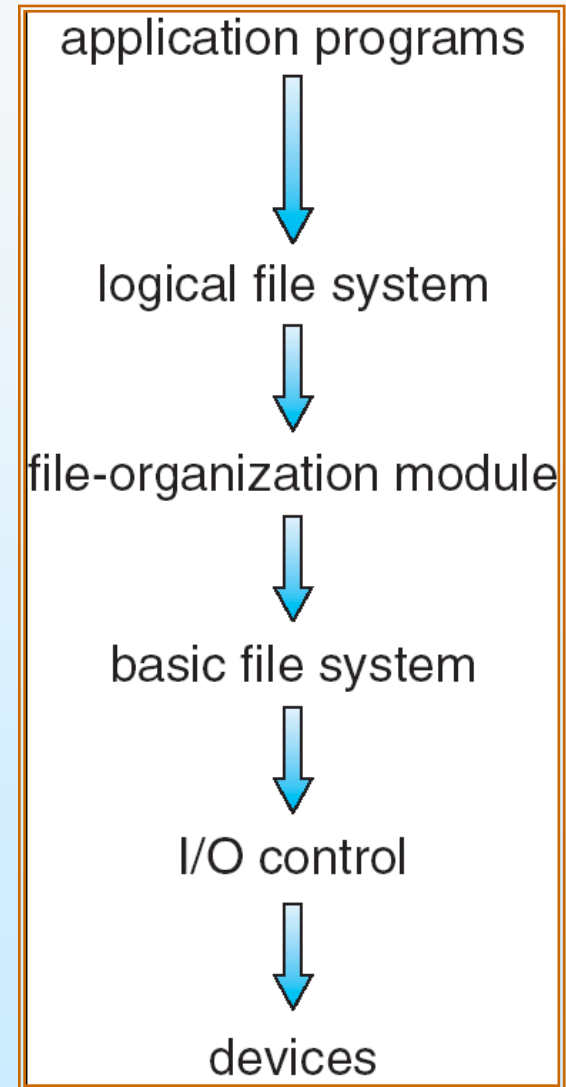
- Ask manager to create a group (unique name), say G, and add some users to the group.
- For a particular file (say *game*) or subdirectory, define an appropriate access.



Attach a group to a file

File-System Structure

- File structure
 - Logical storage unit
 - Collection of related information
- File system resides on secondary storage (disks)
- File system organized into layers



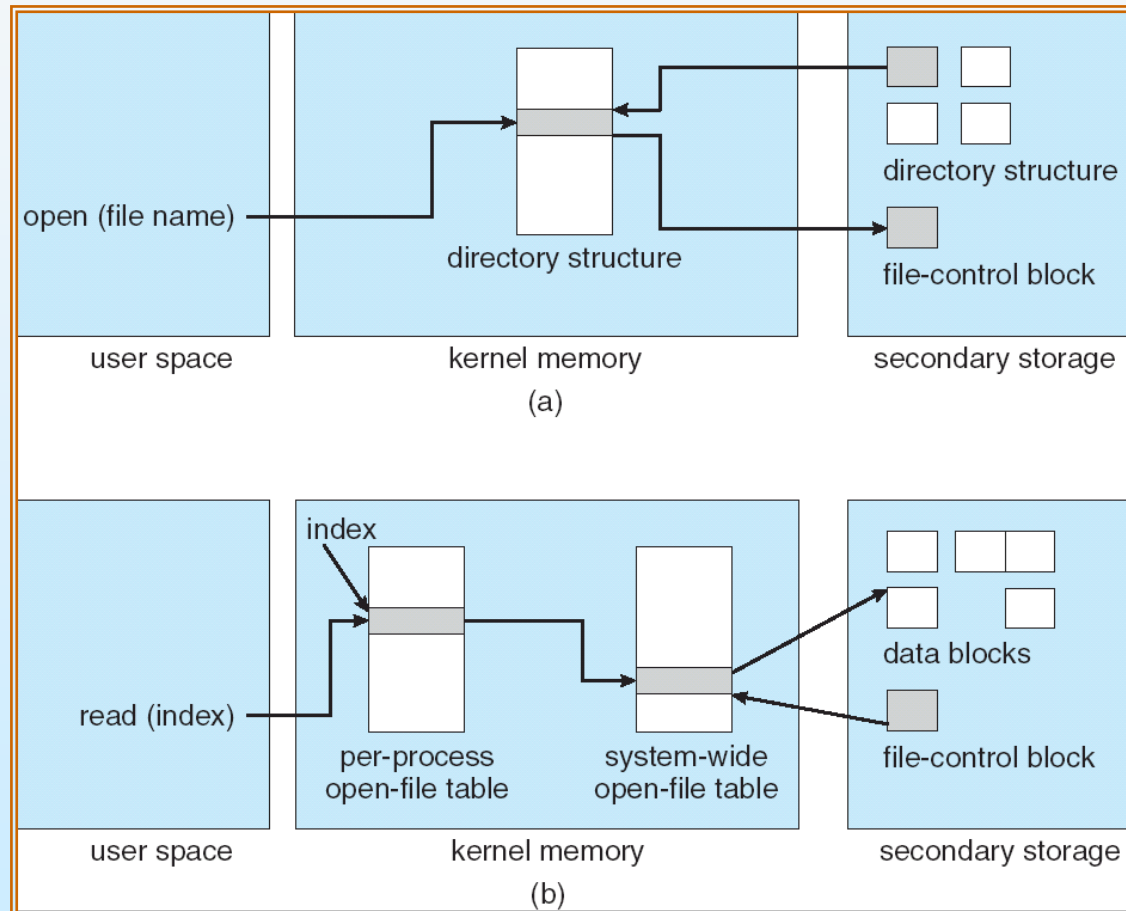
File-System Structure - FCB

- **File control block** – storage structure consisting of information about a file

file permissions
file dates (create, access, write)
file owner, group, ACL
file size
file data blocks or pointers to file data blocks

In-Memory File System Structures

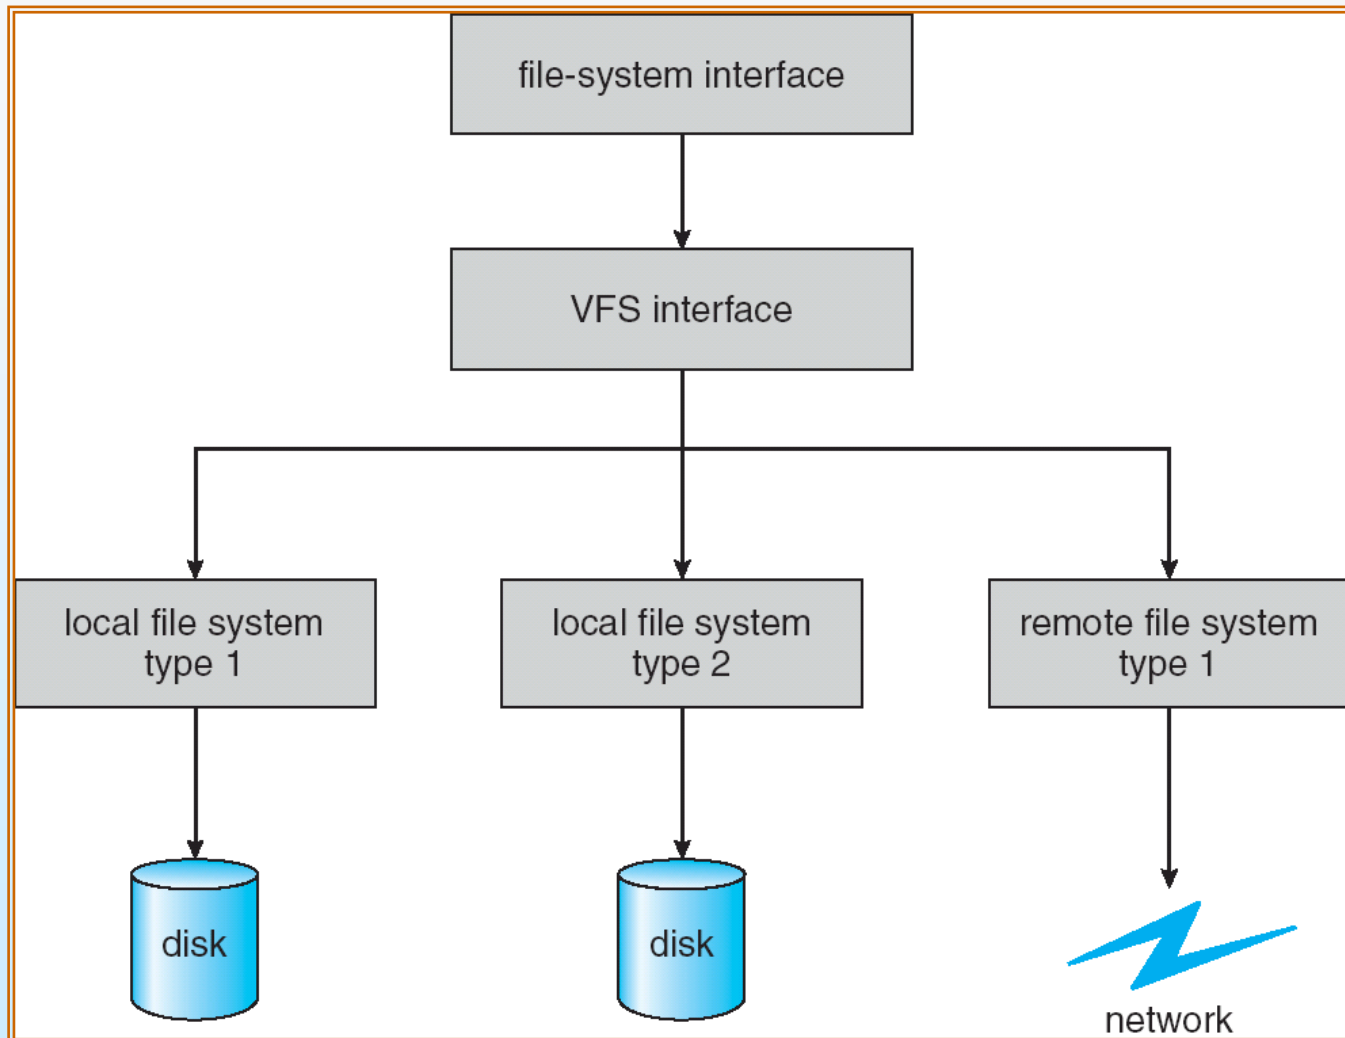
- The following figures illustrates the necessary file system structures provided by the operating systems.



Virtual File Systems

- Virtual File Systems (VFS) provide an object-oriented way of implementing file systems.
- VFS allows the same system call interface (the API) to be used for different types of file systems.
- The API is to the VFS interface, rather than any specific type of file system.

Schematic View of Virtual File System



Directory Implementation

- **Linear list** of file names with pointer to the data blocks.
 - simple to program
 - time-consuming to execute
 - ▶ Create vs Delete
- **Hash Table** – linear list with hash data structure.
 - decreases directory search time
 - **collisions** – situations where two file names hash to the same location
 - fixed size
- **Other data structures**
 - Sorted list, B-Tree, etc.

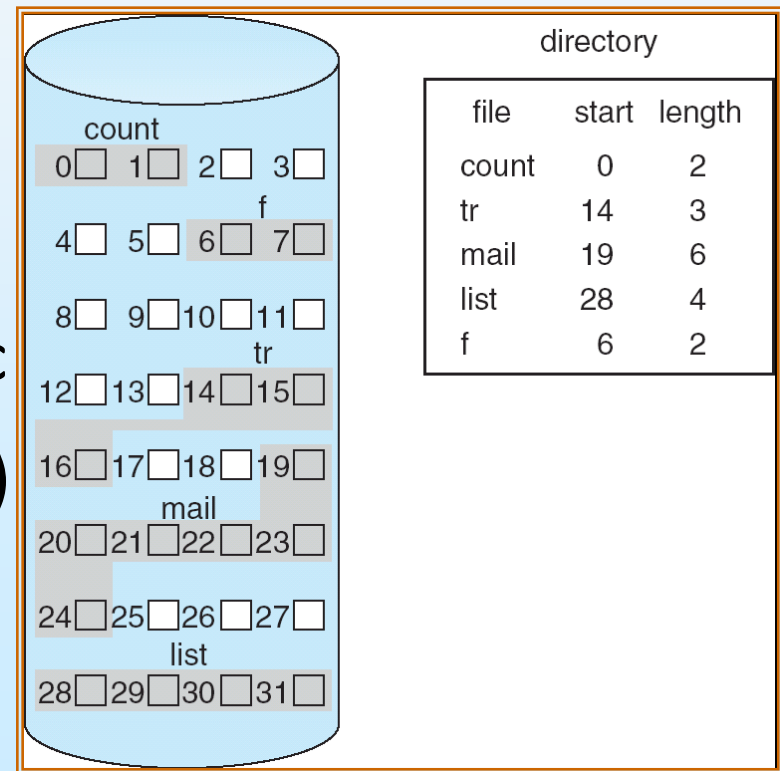
Disk Allocation Methods

■ An allocation method refers to how disk blocks are allocated for files:

- **Contiguous allocation**
- **Linked allocation**
- **Indexed allocation**


Contiguous Allocation

- Each file occupies a set of contiguous blocks on the disk
- Simple – only starting location (block #) and length (number of blocks) are required
- Random access
 - Sequential access ($b+1$)
 - Direct access ($b+i$)
- Wasteful of space (dynamic storage-allocation problem)
- Files cannot grow easily




Linked Allocation

- Each file is a linked list of disk blocks: blocks may be scattered anywhere on the disk.



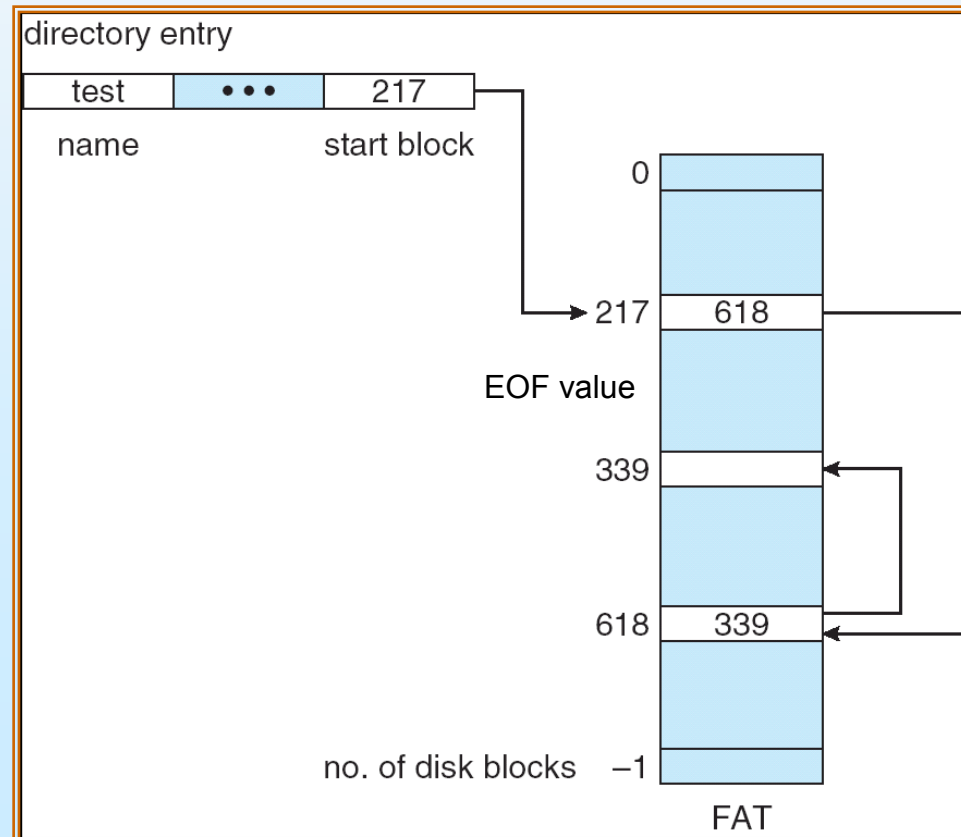
Linked Allocation (Cont.)

- Simple – need only starting address
- Free-space management system
 - No waste of space
 - Really?
- Cluster based allocation
- No random access
 - Why?
- Reliability Issue



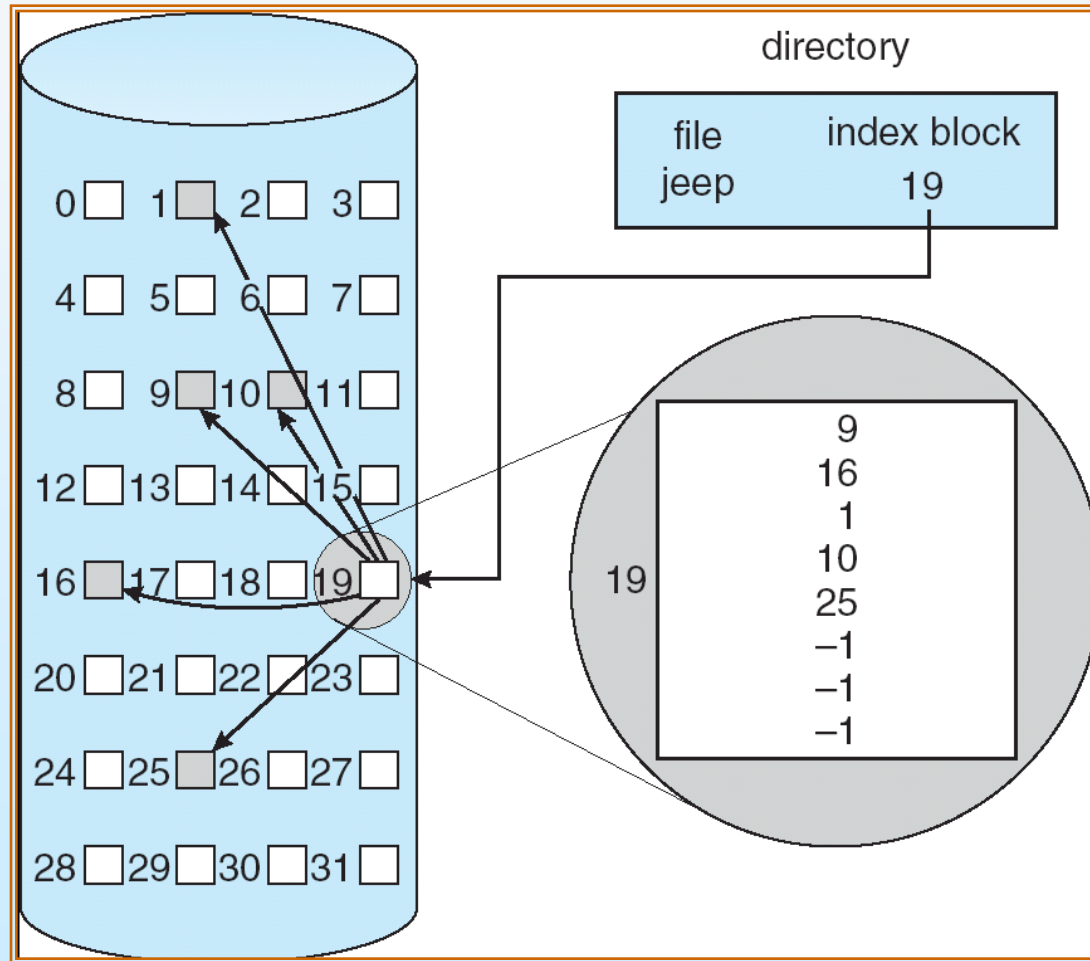
File-Allocation Table

- A variation of the Linked Allocation.
 - A section of disk at the beginning of each volume is set aside.
 - Simple and efficient
 - Entry represents disk block.
 - EOF vs 0 values
 - 0 means unused blocks
- Improved random-access time.
- Disk head seeks increases
 - Why?



Indexed Allocation

- Brings all pointers together into the *index block*.




Indexed Allocation (Cont.)

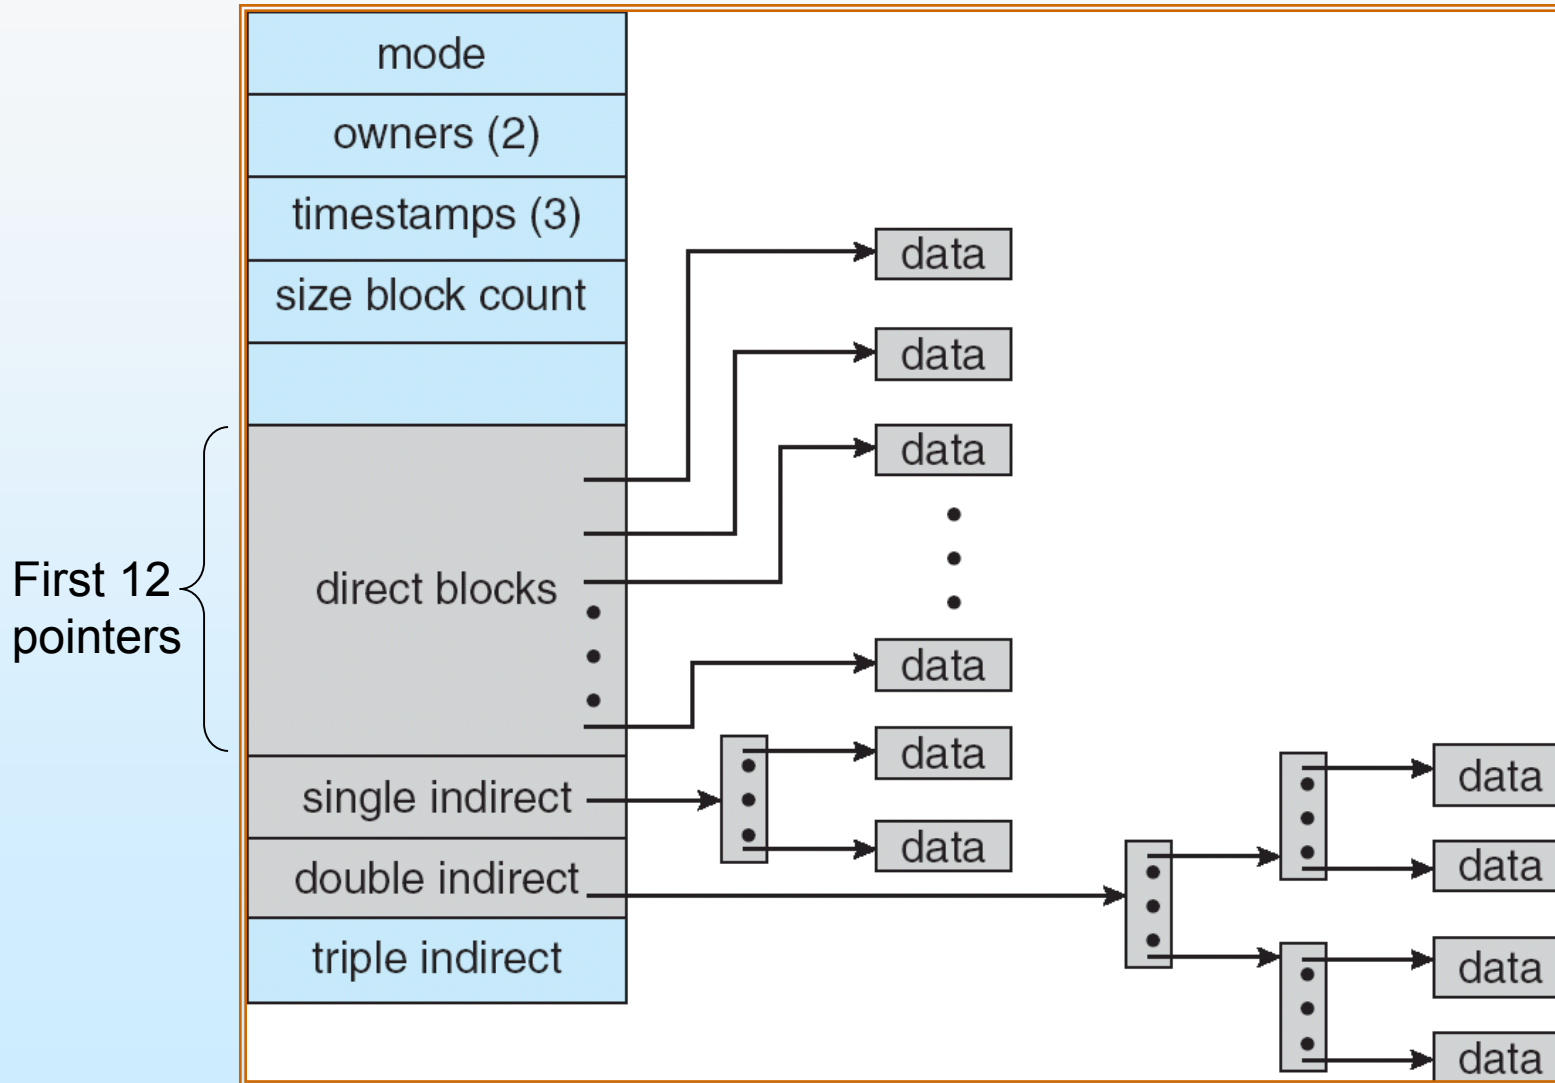
- Need index table
- Random access
- Dynamic access without external fragmentation, but have overhead of at least one index block.
- Mapping from logical to physical in a file of maximum size of 256K words and block size of 512 words.
 - How many blocks we need for the index table?
- Mapping from logical to physical in a file of unbounded length (block size of 512 words).

How large the index block should be?

- Linked scheme – Link blocks of index table (no limit on size).
 - Last word of the index block is a *nil* or a pointer.
- Two-level index
 - What is the maximum size of the file? Assume the disk block is 512 Bytes and a pointer needs 4 bytes.



Combined Scheme: UNIX (4KB per block)



Performance Comparison

■ Contiguous Allocation

- One access to get a disk block
- Maximum length has to be determined at the beginning of creation

■ Linked Allocation

- Direct Access needs i disk reads to access to i^{th} block
- It is fine with sequential access

■ Indexed Allocation

- Index block requires considerable space
- Index structure, file size determine the performance

■ Combined approach

- Using contiguous allocation for small file
- Using indexed allocation for large file